

Highlights

Conformal Surprisal Fusion for Detecting Prompt-Injection Compromise in LLM Agent Traces

Quang-Vinh Dang

- Fuses observation- and action-side trace signals to detect agent compromise
- Conformal surprisal fusion is scale-free and needs no tuned signal weights
- Attains 0.957 AUROC and 0.919 TPR at a realised 3.4% false-positive rate
- Subtracting the benign maximum provably caps AUROC at one half
- AgentDojo hides almost no compromise: its trace-only ceiling does not bind