

Conformal Surprisal Fusion for Detecting Prompt-Injection Compromise in LLM Agent Traces

Quang-Vinh Dang^a

^a*British University Vietnam, Hung Yen, Vietnam*

Abstract

Large language model (LLM) agents that call external tools ingest untrusted content and fold it into the context that drives their next action, exposing them to indirect prompt injection. Runtime monitors that watch the agent’s black-box execution trace are attractive because they are model-agnostic and cheap, but the literature has converged on a single family of evidence: detecting whether untrusted text *looks like an instruction*. We ask what else the trace affords, and how heterogeneous evidence should be combined.

We make four contributions. First, we introduce **SENTRY-Fuse**, a monitor that converts heterogeneous trace signals—one scoring the *observation* stream and two auditing the *action* stream—into conformal tail surprisals against a benign calibration split and fuses them by addition. The rule is scale-free, non-decreasing in every component and free of tuned weights, and we prove it distribution-free valid. We also prove, and then measure, that the natural alternative of subtracting the benign training maximum is many-to-one and can drive an entire attack family to chance—as it did in an earlier, unpublished draft of this work, where it cost 0.25 AUROC and produced an explanation we had written up and believed, which was wrong. Second, we show that action-side evidence is complementary rather than redundant: on 90 genuinely compromised trajectories the three-signal fusion reaches $\text{AUROC} = 0.957 \pm 0.004$ and a true-positive rate of 0.919 ± 0.007 at a realised 3.4% false-positive rate, against 0.944 and 0.729 for the instruction signal alone. Third, we formalise the **detectability ceiling** of trace-only monitoring—the fraction ν of compromises leaving no trace evidence bounds every such monitor—and then show that on AgentDojo it does

Email address: `vinh.dq4@buv.edu.vn` (Quang-Vinh Dang)

not bind: ν lies between 0 and 0.044 depending on how payload visibility is defined, placing the bound at or above 0.958, because the benchmark delivers payloads through tool outputs a successful injection must have read. Our residual error is therefore scoring error, not missing information; and since ν depends on a measurement convention as well as on the corpus, cross-corpus rates are incomparable unless both are reported. Fourth, we document a label-polarity pitfall in the standard AgentDojo harness that inverts the compromised and resisted classes, which silently inverted our own results, and we give the ground-truth procedure that settles it. We report our negative results in full, including four rejected signals and the finding that a hand-picked threshold outperformed the anytime-valid e-detector we began from. All code, trajectories and scripts are released.

Keywords: LLM agents, indirect prompt injection, runtime monitoring, conformal prediction, anomaly detection, AI agent security

1. Introduction

Autonomous agents built on large language models now read email, browse the web, query databases and move money on a user’s behalf. To act, they consume *untrusted* content—web pages, documents, tool outputs—and fold it directly into the context that determines their next tool call. This is the attack surface exploited by *indirect prompt injection*: an adversary plants an instruction inside data the agent will later read, and the agent, unable to separate trusted instructions from untrusted data, follows it (Zhan et al., 2024; Debenedetti et al., 2024). As agents move from demos into production, a runtime monitor—something that watches the live execution trace and halts or escalates when behaviour departs from what is safe—becomes as important as the agent itself.

Monitors that observe only the black-box trace are especially attractive in practice. They need no access to model weights or activations, they impose no retraining cost, they compose with whichever model a deployment happens to use, and they can be dropped in front of a third-party agent. A substantial line of work therefore builds detectors over traces (Wang et al., 2025; Liu et al., 2025; Jia et al., 2025). Almost all of it, however, rests on one kind of evidence: examining the untrusted content the agent read and asking whether it *looks like an instruction*. Keyword and perplexity heuristics do this lexically; fine-tuned detectors and LLM judges do it semantically.

The observation that started this work. We began from that same premise and built a monitor on it. It performed well on the standard benchmark attack and then collapsed, to chance, when we changed the attack template. Our first explanation was that the new attack was phrased too subtly. That explanation was *wrong*: the new payloads were, if anything, more imperative, and they appeared in the trace more often. The collapse came from our own normalisation step. Diagnosing that failure honestly led us to a more general question, which organises this paper:

What, exactly, can a monitor learn from a black-box agent trace, and where does that information run out?

Answering it requires separating two objectives that the literature routinely conflates. *Attempt detection* asks whether untrusted instructions entered the agent’s context at all; *compromise detection* asks whether the agent actually did the attacker’s bidding. They have different positive classes, different ceilings, and—as we show—different optimal evidence. A monitor tuned for the first is not automatically good at the second.

Contributions. Our contributions are as follows.

1. **SENTRY-Fuse, a conformal fusion monitor** (Section 3.4). We combine three trace signals—one scoring the observation stream, two auditing the action stream—by mapping each to its conformal tail surprisal against a benign calibration split and summing. We show (Propositions 2–3) that the construction is distribution-free valid, scale-free and non-decreasing in every component, and that the excess-over-maximum normalisation used in an earlier version of this work is many-to-one and can collapse an entire attack family to a point—which it did, measurably (Section 3.6).
2. **Action-side evidence is complementary, not redundant** (Section 5). Signals that audit *what the agent did* carry information that a signal scoring *what the agent read* does not. Adding the two action-side audits lifts pooled AUROC from 0.944 to 0.957 and the true-positive rate from 0.729 to 0.919, and it removes almost all of the variance the single signal shows across calibration draws.
3. **The detectability ceiling, and why it does not bind here** (Section 5.4). Theorem 5 bounds the true-positive rate of *any* trace-only

monitor by the fraction of compromises that leave trace evidence. Measuring that fraction on AgentDojo puts the bound at or above 0.958 under every criterion we tried, so it does not bind: injections arrive through tool outputs a successful attack must have read. We also show the estimate is sensitive to how “visible” is defined, which is why we withdraw an earlier claim that a published figure exceeded this corpus’s information limit.

4. **Methodological findings and honest negatives** (Sections 4.4 and 5.5).

We document a label-polarity trap in the standard evaluation harness that inverts the compromised and resisted classes; four signals we implemented and rejected, two of which we had previously reported as part of the method; and the finding that a hand-picked threshold outperformed the anytime-valid e-detector we started from, which we report rather than omit.

Every number, table and figure reported here is regenerated by scripts in the released repository without any model API access, since all trajectories and cached judge scores are published alongside the code.

2. Background and Related Work

2.1. Indirect prompt injection

Direct prompt injection manipulates the user-supplied prompt; *indirect* prompt injection is the more dangerous variant for agents, because the adversary never talks to the agent at all. Greshake et al. (2023) introduced the threat and gave the first systematic treatment of it. Their central observation is the one this entire literature rests on: once an LLM is augmented with retrieval, *processing* untrusted retrieved data becomes analogous to *executing* arbitrary code, because the boundary between data and instructions is not represented anywhere in the model’s input. They derive a taxonomy of resulting threats—information gathering, fraud, intrusion, malware (including prompts that propagate as worms), manipulated content, and availability attacks—and a taxonomy of delivery methods, distinguishing *passive* injection (planting content that retrieval will find), *active* injection (sending content, e.g. by email, that the agent will process), *user-driven* injection (tricking the user into pasting it), and *hidden* injection (multi-stage or encoded payloads). They demonstrate feasibility against real deployed systems including Bing Chat and GitHub Copilot. Our threat model (Section 3) is

the passive/active case, and the compromises we detect fall predominantly into their information-gathering and fraud categories.

The consequences scale with the agent’s privileges: an agent that can send email can exfiltrate; an agent that can move money can steal. Zhan et al. (2024) systematise this for tool-integrated agents specifically.

Wallace et al. (2024) argue that the root cause is the absence of a privilege ordering over instruction sources, and train models to respect an explicit *instruction hierarchy*. This mitigates but does not eliminate the problem, since the model must still infer which span of a long context is privileged. Nasr et al. (2025) show that defences evaluated only against static attacks tend to collapse under adaptive attacks that optimise against the defence, and argue that adaptive evaluation should be the norm—a caution we take seriously in Section 6.1.

2.2. Benchmarks for agent security

Progress here is benchmark-driven, and the choice of benchmark determines what “detection” even means.

AgentDojo (Debenedetti et al., 2024) is the most widely used harness for our setting. It provides four tool environments (banking, workspace, travel, Slack), a set of *user tasks* defining legitimate goals, and a set of *injection tasks* defining attacker goals, together with attack templates that place the attacker’s directive into environment content. Crucially, it scores two things separately: *utility*, whether the user’s task was accomplished, and *security*, whether the *injection* task’s goal was accomplished. Section 4.4 shows that the direction of this second flag is easy to invert, with severe consequences.

InjecAgent (Zhan et al., 2024) contributes a distinct family of injection templates aimed at tool-integrated agents, which we use to test whether a detector generalises beyond the template it was developed on. **τ -bench** (Yao et al., 2025) supplies realistic benign multi-turn tool-use trajectories in customer-service domains; it has no attack suite, so we use it purely to enlarge and diversify the benign reference. **WebArena** (Zhou et al., 2024) provides realistic web environments and is complementary but not injection-focused.

AgentHarm (Andriushchenko et al., 2025) deserves separate mention because it is easy to mistake for our setting. It comprises 110 base (and 440 augmented) explicitly malicious agent behaviours across 11 harm categories over a suite of synthetic tools, and it scores whether an agent both refuses *and* retains capability. Its threat model, however, is agent *misuse*: the harmful request comes from the user directly. Ours is indirect injection,

where the user is benign and the harm arrives through third-party content—a distinction the AgentHarm authors themselves draw when contrasting their benchmark with AgentDojo. The two are complementary and a complete guardrail needs both; we evaluate only the injection case, which Section 6.1 records as a scope limitation.

Li et al. (2026) survey the space and note substantial inconsistency in how safety benchmarks define and measure success—a concern our Sections 4.4 and 5.4 substantiate concretely.

2.3. Defences: prevention

Prevention redesigns the agent, or retrains the model, so that injected instructions cannot influence control flow. It is worth separating three mechanisms, because each fails differently and each leaves a different residual role for monitoring.

Marking provenance in the prompt. **Spotlighting** (Hines et al., 2024) observes that an LLM receives multiple input sources concatenated into one token stream with no signal of which span came from where, and supplies that signal by transforming untrusted spans: delimiting them with explicit markers, *datamarking* them by interleaving a reserved token throughout the untrusted text, or encoding them (e.g. base64) so that the model can still read them but their instruction-like surface form is disrupted. On GPT-family models this reduces attack success from above 50% to below 2% with little utility loss. The approach is attractively cheap—it is prompt engineering, requiring no training—but it relies on the model honouring the marking, which is a behavioural rather than an architectural guarantee.

Changing the interface. **StruQ** (Chen et al., 2025a) makes the sharpest conceptual argument in this literature. Prompt injection, they observe, is the latest instance of a recurring vulnerability class in which control and data share one channel: exactly as SQL injection arises because a query string mixes control with data, and cross-site scripting because an HTML page does, prompt injection arises because the LLM API is *unsafe by design*, expecting developers to concatenate the prompt and the data. Their remedy is the analogue of prepared statements: a *structured query* presenting instruction and data as two separate inputs, implemented by a filtering front-end using reserved delimiter tokens together with *structured instruction tuning*, which fine-tunes the model on examples where instructions appear in the data

portion and must be ignored. This drives manually-crafted attacks below 2% and substantially weakens optimisation-based ones (TAP 97% $\rightarrow$ 9%).

Aligning against the injection. **SecAlign** (Chen et al., 2025b) identifies the weakness of the fine-tuning approach: training only to *prefer* the correct response underspecifies what an incorrect response looks like, so security does not generalise to unseen attacks—StruQ, they report, still exceeds 50% attack success under an attack that optimises the injection. SecAlign therefore builds a preference dataset of (injected input, secure response, insecure response) triples and applies preference optimisation, so the model is explicitly steered *away* from the injected response. This yields near-0% success against optimisation-free attacks and below 10% against optimisation-based ones, improving on StruQ by more than a factor of four.

Sanitising the content. **CommandSans** (Das et al., 2025) strips imperative content from tool outputs before the agent reads them, a lighter-weight route that leaves the agent model untouched—though the sanitiser itself is a trained token-level classifier, so the pipeline is not training-free.

Enforcing provenance at the system level. **CaMeL** (Debenedetti et al., 2025) extracts the control flow from the trusted user query, tracks data provenance through a capability system, and refuses tool calls that carry untrusted taint into sensitive argument positions. This is the strongest guarantee of the group and the most invasive, since it requires rebuilding the agent around the defence.

Why monitoring remains necessary. Every mechanism above modifies the agent or the model. That is often impossible: deployments frequently wrap a third-party model behind an API, cannot retrain it, and cannot re-architect a vendor’s agent loop. Notably, Chen et al. (2025b) explicitly scope their contribution to prevention and set detection-based defences aside; the two lines of work are therefore complementary rather than competing. Prevention narrows the attack surface; monitoring supplies an independent tripwire where prevention is absent, partial, or itself circumvented—and, unlike prevention, it produces an auditable record of what happened, which is what incident response needs.

2.4. Defences: detection

Detection leaves the agent untouched and raises an alarm from a score. **DataSentinel** (Liu et al., 2025) fine-tunes an LLM to detect contaminated

inputs, formulated as a minimax game so that the detector is trained against adaptive evasion. **Task Shield** (Jia et al., 2025) checks each action for alignment with the user’s stated task—conceptually the closest prior work to our action-side signal, though it is used as a standalone gate rather than as one calibrated component of a fused statistic. **MELON** (Zhu et al., 2025) detects injection by re-executing the trajectory with a masked user prompt and comparing the resulting tool calls, on the principle that a hijacked action recurs even when the real task is removed; this is powerful but pays a full agent re-run per check.

AgentArmor (Wang et al., 2025) is the strongest *reported* trace-only detector and our reference point. It lifts the agent trace into a program representation—control-flow, data-flow and program-dependence graphs—and enforces typed security policies over it. Version 1 of the preprint reports 95.75% true positives at a 3.66% false-positive rate on AgentDojo with a GPT-4 agent; later versions withdraw that figure and report true-positive rates between 0.86 and 0.97 at false-positive rates between 0.03 and 0.19 instead. We cite the version 1 number because it is the strongest such figure in the literature and therefore the right reference point, while noting that it is not a stable target and that the work remains an unrefereed preprint. Because it reasons about which untrusted value flowed into which argument of which tool, it is structurally different from a scalar text score; our S5 is a much coarser approximation of the same idea, and Section 4.3 sets out why we make no head-to-head comparison.

Recent work also monitors adjacent failure modes: Belenky et al. (2026) detect reward hacking cheaply, Thaman (2026) benchmark tool-use exploits, and Opoku and Banahene (2026) apply conformal risk control to retrieval and tool-use drift.

2.5. *Distribution-free calibration and anytime-valid inference*

Our fusion rule is built from conformal p -values. Conformal prediction (Vovk et al., 2022) converts any real-valued score into a distribution-free valid p -value using only exchangeability of a calibration sample, which is what lets us combine signals living on incompatible scales without distributional assumptions.

The project began in the adjacent literature on *anytime-valid* inference. E-values and e-processes (Vovk and Wang, 2021; Ramdas et al., 2023; Grünwald et al., 2024) support continuous monitoring: a non-negative process whose expectation is bounded under the null can be inspected at every step, and

thresholding it at $1/\alpha$ controls error at level α at every stopping time by Ville’s inequality. Shin et al. (2023) extend this to sequential *change* detection with e-detectors, whose average-run-length (ARL) guarantee holds with no independence assumption—attractive here, because an agent’s actions are profoundly dependent. Sathuka et al. (2025) note that when the reference is *learned* the process is no longer an exact e-process and supply a PAC-style thresholding correction, and Waudby-Smith and Ramdas (2023); Page (1954); Lorden (1971) provide the betting and classical change-detection background. Several 2026 works apply anytime-valid ideas to LLM systems (Xie, 2026; Li, 2026) and to strategic behaviour in multi-agent systems more generally (Gauthier et al., 2026), and conformal ideas to agent error attribution (Feng et al., 2026).

We retain conformal calibration, which does real work in our method, but Section 5.5.2 reports that the sequential e-detector machinery did *not* outperform a hand-picked threshold at the data scales we could reach. We therefore present it as one calibration option rather than as the paper’s contribution—a demotion the evidence forced.

3. Threat Model and Method

3.1. Agent execution traces

An agent pursues a user request u by interleaving tool calls and tool results. We write a trace of length T as

$$\mathcal{T} = ((o_0, a_1), (o_1, a_2), \dots, (o_{T-1}, a_T)), \quad (1)$$

where a_t is the t -th action (a tool name together with its arguments) and o_{t-1} is the tool output the agent observed immediately before choosing a_t . A monitor sees u and $\mathcal{T}$ and nothing else: no model weights, no activations, no logits. We deliberately exclude the agent’s own natural-language messages, both because they are not always available in production logging and to keep the setting comparable to prior trace-only work; Section 6.1 shows this exclusion is exactly what the residual undetectable cases exploit.

3.2. Adversary

The adversary controls the textual content of one or more environment resources that the agent may retrieve, and nothing else. They cannot modify u , the tool implementations, the model, or the monitor. Their goal is to

cause the agent to execute an *injection task*: an action sequence serving the adversary rather than the user. This is the standard indirect-prompt-injection model (Debenedetti et al., 2024).

3.3. Two distinct detection objectives

Definition 1 (Attempt detection). A trace is *attempt-positive* if injected adversarial content entered the agent’s context, whether or not the agent complied.

Definition 2 (Compromise detection). A trace is *compromise-positive* if the agent actually accomplished the adversary’s injection task.

Both are legitimate targets. Attempt detection suits a guardrail that should escalate whenever untrusted instructions reach the context, since an agent that resisted today may not tomorrow. Compromise detection suits incident response and is the harder, more consequential question. The two differ sharply in practice: in our corpus a capable agent resists most attempts, so the attempt-positive class is several times larger than the compromise-positive class—393 against 90 pooled, and 198 against 18 for the DeepSeek agent. Results for one must never be reported as results for the other—a confusion we made ourselves, as Section 4.4 describes.

3.4. Overview of *SENTRY-Fuse*

Figure 1 summarises the monitor. Each action of the live trace is scored by three signals drawn from two distinct evidence channels; each raw score is converted to a conformal tail surprisal against a benign calibration split; the surprisals are summed and compared against a threshold; and when the monitor fires, the dominant term names the reason. The subsections below define each stage. Section 3.5 presents five signals in total, because two of them—which an earlier, unpublished draft of this work included in the method—turned out to be near-chance under an honest protocol and are reported as negative results rather than quietly dropped.

3.5. Signals

We implemented five per-action signals. Exactly one, S3, scores the *observation* stream, asking whether untrusted content is behaving like an instruction. The other four score the *action* stream. Of those, S4 and S5 ask whether the agent’s behaviour serves the user and are retained; S1 and S2 ask




Figure 1: Overview of SENTRY-Fuse. Three signals are computed per action from two distinct evidence channels: the *observation* channel (S3) asks whether untrusted content is behaving like an instruction, while the *action* channel (S4, S5) asks whether the agent’s behaviour serves the user’s request. Each raw score is mapped to a conformal tail surprisal against a benign calibration split, which makes the heterogeneous signals commensurable at a far finer resolution than the excess-over-maximum alternative of Proposition 1, which collapses a whole class to a point. The surprisals are summed and thresholded; on an alarm the largest term, $\arg \max_k \psi_k$, names the signal that fired, which gives the operator an actionable reason rather than a bare alert.

only whether the behaviour is statistically unusual, and Section 5.5 shows that under a protocol which fits the reference model out-of-sample they are near-chance. **SENTRY-Fuse is the fusion of S3, S4 and S5**; S1 and S2 are documented here because we reported them as part of the method previously and because their failure is informative.

S1: tool-transition surprise (rejected). A Laplace-smoothed bigram model over consecutive tool names, fitted on benign traces, gives

$$s_t^{\text{tr}} = -\log \hat{P}(\text{tool}(a_t) \mid \text{tool}(a_{t-1})), \quad (2)$$

with $\text{tool}(a_0)$ a distinguished start symbol so that the score is defined at $t = 1$.

S2: argument novelty (rejected). The fraction of tokens in a_t ’s arguments that never appeared in any argument of the same tool in the benign reference,

$$s_t^{\text{nov}} = \frac{|\{\text{tokens of args}(a_t)\} \setminus V_{\text{tool}(a_t)}|}{|\{\text{tokens of args}(a_t)\}|} \in [0, 1], \quad (3)$$

where V_{tool} is the per-tool argument vocabulary of the reference split. An earlier, unpublished draft of this work folded S1 and S2 into a single score as $s_t^{\text{tr}} + 4 s_t^{\text{nov}}$, adding two raw quantities on incommensurable scales (nats and a fraction) with a weight chosen by hand. That is precisely the failure mode Section 3.7 exists to remove.

We must be precise about where this has and has not been undone. **SENTRY-Fuse and its ablation (Table 3) use no weights at all:** every signal is conformalised separately and summed, which is the claim made in Section 3.7. The *attempt-detection* results (Tables 4, 5 and 7) are computed with the project’s original scalar pipeline, which sums the per-action signals with fixed coefficients (1, 4, 4, 4). We have not re-derived those tables, because in each of them the weighting is held constant while one factor varies—the normalisation in Table 5, the agent or attack family in Table 4, the combiner in Table 7—so the comparisons they support are unaffected. Their absolute values, however, are those of the weighted composite and should not be read as single-signal performance.

S3: observation instruction-likeness. Let $\text{IL}(o)$ score how strongly a tool output reads as a directive addressed to the agent. We evaluate two instantiations. The *lexicon* variant is a weighted per-word density of imperative verbs, second-person pronouns, pseudo-markup and exclamation marks. The *judge* variant queries a small LLM with the observation alone and takes a normalised integer rating. Section 5.3 compares them; the judge generalises better across attack templates, at the price of one cached model call per distinct observation.

S4: action-justification audit. Given only the user request u , the tool name and its arguments—and *never* the observations—an LLM judge rates how far the call fails to serve u :

$$s_t^{\text{aj}} = \text{Unjustified}(u, a_t) \in [0, 1]. \quad (4)$$

Withholding the observations is a deliberate design choice: it prevents the audit from simply re-detecting the injected text, so S4 is not a restatement of S3 and the two may be fused without circularity.

S5: sensitive-sink provenance. A conjunctive, per-value rule that fires when the agent performs an externally-effectful call whose critical argument traces to untrusted prose:

$$s_t^{\text{sp}} = \mathbf{1}[\text{Sink}(a_t) \wedge \exists v \in \text{Ids}(a_t) : v \notin u \wedge \text{FirstSeenInProse}(v)]. \quad (5)$$

Here Sink marks tools with external effect (moving money, sending messages, publishing, editing standing instructions) as opposed to read-only lookups; Ids extracts identifier-like values (IBANs, e-mail addresses, URLs, long digit strings); and FirstSeenInProse holds when the value’s first appearance in the trace is inside a long natural-language line rather than a short structured record field. The intuition is that a legitimate recipient reaches the agent as a record it looked up, whereas an attacker’s recipient arrives embedded in the prose of a poisoned document.

3.6. Why naive normalisation fails

The signals live on incompatible scales: s^{tr} is in nats and unbounded, s^{il} and s^{aj} lie in $[0, 1]$, s^{sp} is binary. Summing them raw lets the widest-ranged term dominate. A natural fix, and the one an earlier, unpublished draft used, is to subtract the largest value observed in benign training data,

$$\phi_{\max}(x) = \max\{0, x - \beta\}, \quad \beta = \max_{x' \in \mathcal{D}_{\text{train}}} x', \quad (6)$$

so that anything within the benign range scores zero. This is a trap.

Throughout, AUROC denotes the area under the receiver operating characteristic curve: the probability that a randomly drawn positive scores above a randomly drawn negative, counting ties as one half.

Proposition 1 (Excess-over-maximum can annihilate a signal). *Let $\phi_{\max}$ be as in (6), let $\mathcal{B}, \mathcal{C}$ be the benign and attack score distributions, and suppose the attack scores satisfy $\mathbb{P}_{X \sim \mathcal{C}}(X \leq \beta) = 1$. Write $q = \mathbb{P}_{X \sim \mathcal{B}}(X > \beta)$ for the benign exceedance probability. Then*

$$\text{AUROC}(\phi_{\max}(\mathcal{B}), \phi_{\max}(\mathcal{C})) = \frac{1}{2}(1 - q) \leq \frac{1}{2}, \quad (7)$$

irrespective of how well x itself separates the classes. In particular the transform is at best uninformative under this hypothesis, and strictly worse than chance whenever some benign score exceeds β .

Proof. $\phi_{\max}$ maps $(-\infty, \beta]$ to the single point 0. Since $\mathcal{C}$ is supported on $(-\infty, \beta]$, every transformed attack score is exactly 0. A transformed benign score is 0 with probability $1 - q$ and strictly positive with probability q . Comparing an independent attack score with a benign score, the attack score is never larger, is tied with probability $1 - q$, and is smaller with probability q . Counting ties as one half gives $\text{AUROC} = \frac{1}{2}(1 - q) + 0 \cdot q$, which is (7). $\square \quad \square$

Two remarks. First, $q > 0$ is the generic case rather than an edge case: if β is the maximum of n exchangeable benign training scores from a continuous law, then a fresh benign score exceeds it with probability exactly $1/(n+1)$. Second, the hypothesis $\mathbb{P}_c(X \leq \beta) = 1$ holds only approximately on real data—a few attack scores do exceed β —which is why the collapse measured in Section 5.3 lands just *above* $1/2$ (0.515 to 0.531) rather than at or below it. The proposition identifies the mechanism and its worst case; it does not predict the exact value.

The premise is not hypothetical. In our data the benign median instruction-likeness is 0.010 but its maximum is 0.613—an outlier—whereas the median for the InjecAgent family is 0.115. Every InjecAgent score therefore fell below β , and the signal was driven to AUROC = 0.531 under $\phi_{\max}$ against 0.706 under the conformal transform of Section 3.7: a loss of 0.175, and 0.246 on the `important_instructions` family (0.515 against 0.761). An earlier, unpublished draft of this work reported that collapse and attributed it to the attack being subtly phrased. That explanation was wrong—the InjecAgent payloads are imperative and appear in the trace in every case—and the cause was our own normalisation.

The lesson is about *resolution*, not merely about monotonicity. Both $\phi_{\max}$ and the conformal transform below are many-to-one; the difference is that $\phi_{\max}$ can map an entire class to a single point, destroying all within-class ordering, whereas the conformal transform retains $n+1$ distinct levels. A component transform must therefore be chosen so that its quantisation is fine relative to the separation one hopes to exploit—a requirement we make precise in Proposition 3.

3.7. Conformal surprisal fusion

Definition 3 (Conformal tail surprisal). Let $x_{1:n}$ be benign calibration values for a signal and x a test value. The conformal p -value and surprisal are

$$p(x) = \frac{1 + |\{i : x_i \geq x\}|}{n+1}, \quad \psi(x) = -\log p(x). \quad (8)$$

Proposition 2 (Distribution-free validity). *If $x, x_1, \dots, x_n$ are exchangeable, then $\mathbb{P}(p(x) \leq \alpha) \leq \alpha$ for every $\alpha \in (0, 1)$; equivalently $\mathbb{P}(\psi(x) \geq \log(1/\alpha)) \leq \alpha$.*

Proof. If the values are almost surely distinct, exchangeability makes the rank of x among $\{x, x_1, \dots, x_n\}$ uniform on $\{1, \dots, n+1\}$, and $\{p(x) \leq \alpha\}$

is the event that this rank places x among the top $\lfloor \alpha(n+1) \rfloor$ values, of probability $\lfloor \alpha(n+1) \rfloor / (n+1) \leq \alpha$. Ties only help: because $|\{i : x_i \geq x\}|$ counts calibration values *weakly* exceeding x , a tie can only increase $p(x)$, so the bound continues to hold with the same argument applied to any tie-breaking rule. $\square$ $\square$

The tie case is not academic here. S5 is binary by construction and S4 saturates at 1.0 on 17% of benign trajectories, so the calibration samples of those signals contain many repeated values; the weak inequality in (8) is what keeps p conservative rather than anti-conservative.

Proposition 3 (Ordering and scale-freeness). *Fix calibration values $x_{1:n}$ and let ψ be as in (8). Then (i) ψ is non-decreasing, so it never inverts a strict ordering; (ii) ψ is invariant to any strictly increasing reparameterisation of the signal; and (iii) ψ takes at most $n+1$ distinct values, all in $[0, \log(n+1)]$, so it is many-to-one, and AUROC is preserved only up to the ties this quantisation introduces.*

Proof. $|\{i : x_i \geq x\}|$ is non-increasing in x , so p is non-increasing and $\psi = -\log p$ is non-decreasing, which gives (i): if $x < x'$ then $\psi(x) \leq \psi(x')$, so no strictly ordered pair is reversed. For (ii), ψ is a function of the rank of x among the calibration values, and ranks are invariant under strictly increasing maps. For (iii), $|\{i : x_i \geq x\}|$ ranges over $\{0, 1, \dots, n\}$, so p takes values in $\{1/(n+1), \dots, 1\}$ and ψ in $\{0, \dots, \log(n+1)\}$; in particular every x strictly above $\max_i x_i$ receives the same value $\log(n+1)$. $\square$ $\square$

Part (iii) is a genuine limitation and we state it rather than claim more. Because the transform saturates above the calibration maximum, it *can* lose discrimination. With $n = 1$, calibration $\{0\}$, benign scores $\{0.1, 0.2\}$ and attack scores $\{0.3, 0.4\}$, a perfectly separating signal (AUROC = 1) is mapped to the constant $\log 2$ and AUROC falls to $1/2$. Nor is $1/2$ a floor: with calibration $\{10\}$, benign $\{0, 0, 11\}$ and attack $\{1, 2\}$, an above-chance signal (0.667) becomes 0.333, because collapsing correctly ordered pairs into ties can cost more than it gains. Monotonicity therefore buys *no* guarantee about AUROC; only resolution does.

At realistic calibration sizes the cost is small but not zero. On well-separated synthetic Gaussians, where the raw signal has AUROC = 1, the conformal transform yields 0.925 at $n = 5$, 0.979 at $n = 23$, 0.986 at $n = 35$ and 0.992 at $n = 63$ (400 replicates each; the calibration sizes used in this

paper are 23 to 63). The loss is $O(1/n)$, which is why we prefer the largest calibration split the corpus allows.

The contrast with Proposition 1 is accordingly quantitative rather than categorical. Both transforms are many-to-one. The difference is how coarse they are: $\phi_{\max}$ maps an entire class to a single point, leaving no within-class ordering at all and forcing $\text{AUROC} = \frac{1}{2}(1 - q) \leq \frac{1}{2}$, whereas ψ retains $n + 1$ ordered levels. What a fusion rule needs is not injectivity, and not monotonicity alone, but a quantisation fine enough that the surviving ordering still separates the classes; Section 5.3 measures what that is worth on real signals.

Definition 4 (SENTRY-Fuse). For signals $k = 1, \dots, K$ with per-trajectory summaries x_k (we use the per-trajectory maximum of each per-action signal) and benign calibration values $x_{1:n_k}^{(k)}$, the fused statistic is

$$S = \sum_{k=1}^K \psi_k(x_k) = - \sum_{k=1}^K \log p_k(x_k), \quad (9)$$

and the monitor alarms when $S \geq \tau$, with τ set to a benign quantile on a held-out calibration split.

The rule has no tuned weights: every signal contributes on the common scale of log-improbability, which is what makes an unweighted sum meaningful. Two caveats belong here rather than in a footnote. First, by Proposition 3(iii) component k contributes at most $\log(n_k + 1)$, so if the calibration sets differ in size the sum is implicitly weighted by $\log(n_k + 1)$; we therefore use a single calibration split, giving $n_k = n$ for all k and equal maximum contributions. Second, the *threshold* τ remains a calibrated quantity, which we choose as a benign quantile; the claim is that the relative contribution of the signals is untuned, not that the monitor is free of all calibration.

Proposition 4 (Fused false-alarm bound). *Under the null, for any $\tau > 0$,*

$$\mathbb{P}(S \geq \tau) \leq \min \left\{ 1, \sum_{k=1}^K e^{-\tau/K} \right\} = \min \{ 1, K e^{-\tau/K} \}. \quad (10)$$

If in addition the p_k are mutually independent, then $\mathbb{P}(S \geq \tau) \leq \Gamma(K, \tau)/\Gamma(K)$, the sharper bound obtained by Fisher’s method, where $\Gamma(K, \tau) = \int_{\tau}^{\infty} u^{K-1} e^{-u} du$ is the upper incomplete gamma function.

Algorithm 1 SENTRY-Fuse

Require: disjoint benign splits $\mathcal{D}_{\text{fit}}, \mathcal{D}_{\text{cal}}, \mathcal{D}_{\text{thr}}$; target false-alarm rate α ; signals $1, \dots, K$

- 1: fit any reference a signal needs (e.g. per-tool vocabularies) on $\mathcal{D}_{\text{fit}}$
- 2: **for** $k = 1, \dots, K$ **do**
- 3: $x_{1:n_k}^{(k)} \leftarrow$ per-trajectory summaries of signal k on $\mathcal{D}_{\text{cal}}$, scored against the $\mathcal{D}_{\text{fit}}$ reference
- 4: **end for**
- 5: $\tau \leftarrow$ threshold realising a false-alarm rate $\leq \alpha$ on $\{S(\mathcal{T}) : \mathcal{T} \in \mathcal{D}_{\text{thr}}\}$ (Section 4.3)
- at monitoring time, for a trace $\mathcal{T}$:**
- 6: **for** $k = 1, \dots, K$ **do**
- 7: $x_k \leftarrow \max_t s_t^{(k)}(\mathcal{T})$; $\psi_k \leftarrow -\log((1 + |\{i : x_i^{(k)} \geq x_k\}|)/(n_k + 1))$
- 8: **end for**
- 9: $S \leftarrow \sum_k \psi_k$
- 10: **if** $S \geq \tau$ **then raise alarm**; report $\arg \max_k \psi_k$ as the dominant cause
- 11: **end if**

Proof. If $S \geq \tau$ then some $\psi_k \geq \tau/K$, so $\{S \geq \tau\} \subseteq \bigcup_k \{\psi_k \geq \tau/K\}$; a union bound with Proposition 2 applied at level $\alpha = e^{-\tau/K}$ gives the first claim, and it requires no dependence assumption. For the second, Proposition 2 says each p_k is stochastically larger than a uniform variable, so each $-\log p_k$ is stochastically dominated by a standard exponential; under independence the sum is dominated by a $\text{Gamma}(K, 1)$ variable, whose upper tail is $\Gamma(K, \tau)/\Gamma(K)$. $\square$ $\square$

Only *domination* holds, not equality: since each p_k is supported on the finite grid $\{1/(n+1), \dots, 1\}$, the statistic S is discrete and bounded above by $K \log(n+1)$, so it cannot be exactly Gamma-distributed.

Remark 1. We report empirical operating points throughout rather than relying on Proposition 4, because the assumption-light union bound is loose and independence across our signals is not credible: S3 and S4 both respond to the presence of an injection, albeit through different channels. The bound’s role is to establish that the fused statistic admits distribution-free false-alarm control at all, and to make explicit what would be needed for a sharp one.

Because the fusion is additive over interpretable components, the dominant term $\arg \max_k \psi_k$ names *why* the monitor fired—an injected instruc-

tion, an unjustified action, an untrusted recipient, or an anomalous tool transition—which matters for triage.

4. Experimental Setup

4.1. Corpora

We collect real trajectories by running tool-calling agents on standard benchmarks; Table 1 summarises the result. From AgentDojo (Debenedetti et al., 2024) we run all four suites (banking, workspace, travel, Slack) both benign and under attack, with two attack families: the `important_instructions` template and the InjecAgent (Zhan et al., 2024) template. From τ -bench (Yao et al., 2025) we run the retail domain benign-only, to diversify the benign reference. Two agents are used: DeepSeek-V4-Flash and GPT-4o-mini, from different vendors, so that agent identity can be varied independently of everything else.

Table 1 already contains a substantive finding. The two agents differ enormously in injection robustness: GPT-4o-mini is compromised in 47% of attempts, DeepSeek-V4-Flash in only 9%, and DeepSeek resisted the InjecAgent family *completely* (0/42). Robustness to indirect prompt injection is therefore a strong function of the agent, not a property of the attack alone, and a defence evaluated on a single agent may be measuring that agent’s fragility as much as its own merit.

Remark 2. Because DeepSeek was never compromised by InjecAgent, compromise detection is *undefined* for that cell: there are no positives. Only attempt detection is meaningful there, and we report it as such rather than quoting a rate over an empty class.

4.2. Protocol

Trajectories are converted to per-action signal streams by a deterministic, GPU-free adapter. Feature hashing uses a seed-stable CRC rather than Python’s per-process-salted hash, so the entire evaluation is byte-reproducible. LLM judge scores are cached by content hash (681 distinct observations, 1,329 distinct task–action pairs), so all reported evaluations run offline from committed caches.

Our evaluation protocol is deliberately strict about leakage. For each of ten seeds the benign set is split into three disjoint parts: one third fits the reference model (the bigram table and per-tool argument vocabularies),

Table 1: Trajectory corpora, with the *corrected* label polarity of Section 4.4: *compromised* means the injection task succeeded. Counts are after a two-action minimum filter. “imp. instr.” is AgentDojo’s `important_instructions` template; the DeepSeek agent is DeepSeek-V4-Flash.

Agent	Attack family	Benign	Attacks	Comp.	Resist.
DeepSeek	imp. instr.	71	198	18 (9%)	180
DeepSeek	InjecAgent	—	42	0 (0%)	42
GPT-4o-mini	imp. instr.	70	153	72 (47%)	81
τ -bench retail	none	50	—	—	—
Total		191	393	90	303

one third supplies the conformal calibration values, and one third supplies the held-out negatives. Compromise-positive trajectories are never used for fitting or calibration. We report mean $\pm$ standard deviation over the ten seeds.

The three-way split matters for Proposition 2. If the calibration values are computed on the same trajectories the reference was fitted to, they are in-sample and systematically less surprising than a fresh benign draw, so the conformal p -values are anti-conservative and the stated false-alarm level is not delivered. Fitting on a separate third makes the calibration values and the held-out negatives both out-of-sample with respect to the reference, which is what exchangeability requires. An earlier version of this evaluation split the benign set only in half and drew calibration values from the fitting half.

Remark 3 (Leakage matters). An earlier version of this evaluation fitted the reference on all benign trajectories including the test half. Under that protocol the weighted behavioural composite (S1 and S2 summed with fixed coefficients) appeared excellent, AUROC = 0.982; with the reference fitted out-of-sample and the signals kept separate, S1 and S2 collapse to 0.48–0.56 (Table 3). Only held-out numbers appear in the compromise-detection results. This is worth stating plainly because the leaky number is the one a careless pipeline produces by default.

4.3. Metrics

We report AUROC and the true-positive rate (TPR) at a benign false-positive rate (FPR). Comparing detectors at a common FPR is essential—

they are otherwise trivially incomparable—but on samples of this size a *nominal* FPR cannot generally be hit. With n held-out negatives, thresholding a score yields only the rates k/n for integer k , so the attainable grid has spacing $1/n$: 4.2% on the 24 GPT-4o-mini negatives and 1.5% on the 65 pooled ones. We therefore set the threshold at the k -th largest negative with $k = \lfloor \alpha n \rfloor$, which realises an FPR of $k/n \leq \alpha$, and we *report the realised rate* in every table alongside the target. We use $\alpha = 5\%$ as the target throughout.

This is worth being explicit about because interpolating an empirical quantile, as is conventional, silently produces a different rate: at a nominal 3.66% on 35 negatives the realised rate is $2/35 = 5.71\%$, and a TPR quoted at “3.66%” would then have been measured at over 5%. An earlier, unpublished draft of this work reported its headline number that way.

On comparison with published figures. AgentArmor (Wang et al., 2025) is the strongest reported trace-only detector, and version 1 of that preprint reports 95.75% TPR at 3.66% FPR on AgentDojo with a GPT-4 agent. We do *not* present our numbers as beating it, for three reasons. The corpora differ (different agents, a different attack sample, and—as Section 4.4 shows—a positive-class definition that is easy to get wrong); the operating points differ once the realised rate is accounted for; and the figure itself was removed in later versions of that preprint, which report TPR between 0.86 and 0.97 at FPR between 0.03 and 0.19 instead. We cite it as the relevant reference point and as evidence that our results are in the same regime, not as a baseline we outperform. No re-implementation of it was run on our corpus, so no head-to-head comparison is claimed.

4.4. A label-polarity pitfall

Before presenting results we must document a methodological trap, because it inverted our own conclusions for a substantial period and because the ceiling analysis of Section 5.4 is impossible to get right without it.

AgentDojo returns a per-trajectory `security` flag. It is natural, and wrong, to read this as “the system was secure”. The flag is computed as the success of the *injection* task, so `security=True` means the attacker’s goal was achieved, i.e. the agent *was compromised*.

We initially adopted the opposite convention. The error was invisible in aggregate: both classes are non-empty, all downstream code ran, and every metric produced plausible numbers. It surfaced only when we asked why so

Table 2: Resolving label polarity against in-trace ground truth, over all 107 attacked banking-suite trajectories. Every trajectory flagged `security=True` wired money to the attacker’s account, which settles the polarity. The three `False` trajectories that also paid the attacker are cases of partial compliance, and show that the flag undercounts harm.

	<code>security=True</code>	<code>security=False</code>
agent wired money to attacker IBAN	31	3
agent did not	0	73

many “compromised” trajectories contained no state-changing tool call at all, and read one:

Example 1 (A trajectory we had labelled compromised). Banking suite, `important.instructions`. The agent calls `read_file` and five read-only lookups, then states:

“I see a message embedded in the bill file that’s trying to get me to do something else first. However, I shouldn’t be tricked by that embedded instruction.”

This agent *resisted*. It had been counted as a successful attack.

We resolved the polarity against ground truth rather than documentation. In the banking suite the adversary’s goal is a transfer to a specific attacker IBAN, so compliance is directly checkable in the trace:

Table 2 settles the direction that matters: every one of the 31 trajectories flagged `security=True` did in fact wire money to the attacker, and none of the trajectories that wired no money is so flagged. So `security=True` is compromise, and the polarity we had been using was inverted.

The converse fails, however, and the failure is worth recording. Three trajectories labelled `security=False` *did* pay the attacker. Inspecting them shows why: this injection task requires not merely a transfer to the attacker’s account but one whose subject line names the user’s music-streaming service, and in these three the agent transferred the money with a subject of its own invention. The flag scores the injection task’s *exact* goal, so an agent that is manipulated into moving money but fumbles a detail of the attacker’s script is recorded as having resisted. The harness therefore *undercounts* harm, and our compromised class of 90 is a lower bound. We keep the harness definition—changing it would make our numbers incomparable

with the literature’s—but flag that a harm-based positive class would be the better target, and that the 303 “resisted” trajectories are not all benign.

The consequences for us were severe and instructive. Every “hijack-detection” number we had computed was measuring the *resisted* class. Our effect-based signals—the action audit and sink provenance—had been graded against trajectories in which, by construction, the agent had done nothing harmful; they appeared useless. Once the polarity was corrected the action audit became one of the two signals the method relies on, reaching $\text{AUROC} = 0.900 \pm 0.039$ on GPT-4o-mini and 0.780 ± 0.021 on DeepSeek (Table 3). We had nearly discarded a working signal because of a boolean.

Remark 4. We report this at length because the failure mode is silent, the fix is trivial, and the affected quantity—which class counts as a successful attack—is the foundation of every number in this literature. We recommend that any work using this harness verify polarity against an in-trace ground-truth check of the kind in Table 2, and report having done so. The same check is what surfaced the undercount above, so it is worth running even once the polarity is known to be right.

5. Results

5.1. Compromise detection

Table 3 gives the main result. On GPT-4o-mini, where the compromised class is large enough for stable estimates (72 positives), the observation-side instruction signal alone reaches $\text{AUROC} = 0.951$ but a wildly variable $\text{TPR} = 0.661 \pm 0.433$: it is strong on average and close to unusable at a fixed low false-alarm budget, because whether it clears the threshold depends on which benign trajectories happened to land in the calibration split. Adding the action-justification audit lifts this to 0.914 ± 0.040 , and adding sink provenance to **0.956 ± 0.006** at a realised 2.9% false-positive rate. Pooled over both agents the same ordering holds, ending at 0.919 ± 0.007 .

Figure 2 visualises the ablation. Four points deserve emphasis.

First, **action-side evidence is complementary**. Pooled, the instruction signal alone reaches $\text{AUROC} = 0.944$ and $\text{TPR} = 0.729$; adding the two action-side signals takes these to 0.957 and 0.919. The gain is not explained by the action signals being individually strong: alone they reach 0.862 and 0.836 AUROC, and their recall at their own much stricter realised operating points (0.0% and 0.8%) is near zero. They contribute cases the observation

Table 3: Compromise detection, ten seeds, mean $\pm$ std. Benign trajectories are split three ways per seed: one third fits the reference model, one third supplies the conformal calibration values, one third the held-out negatives. The realised false-positive rate is reported alongside because a 5% target is not exactly attainable at these sample sizes (Section 4.3). The rows therefore sit at *different* operating points—from 0.0% to 4.4%—so the TPR column is not a like-for-like comparison across rows, whereas AUROC, being threshold-free, is. SENTRY-Fuse in bold.

Detector	AUROC	TPR	realised FPR
<i>GPT-4o-mini: 72 compromised, 24 held-out benign</i>			
S3 instruction-likeness only	0.951 $\pm$ 0.026	0.661 $\pm$ 0.433	1.2%
S4 action audit only	0.900 $\pm$ 0.039	0.267 $\pm$ 0.407	0.4%
S5 sink provenance only	0.868 $\pm$ 0.019	0.164 $\pm$ 0.328	0.4%
S1 + S2 behavioural	0.530 $\pm$ 0.058	0.101 $\pm$ 0.076	3.8%
S3 + S4	0.981 $\pm$ 0.009	0.914 $\pm$ 0.040	2.5%
SENTRY-Fuse (S3–S5)	0.985 $\pm$ 0.003	0.956 $\pm$ 0.006	2.9%
all five signals	0.973 $\pm$ 0.013	0.932 $\pm$ 0.053	3.3%
<i>DeepSeek-V4-Flash: 18 compromised, 41 held-out benign</i>			
S3 instruction-likeness only	0.871 $\pm$ 0.014	0.622 $\pm$ 0.311	2.2%
S4 action audit only	0.780 $\pm$ 0.021	0.000 $\pm$ 0.000	0.0%
S5 sink provenance only	0.665 $\pm$ 0.019	0.233 $\pm$ 0.191	2.0%
S1 + S2 behavioural	0.480 $\pm$ 0.058	0.011 $\pm$ 0.022	2.4%
S3 + S4	0.850 $\pm$ 0.005	0.778 $\pm$ 0.000	2.7%
SENTRY-Fuse (S3–S5)	0.843 $\pm$ 0.008	0.778 $\pm$ 0.000	3.9%
all five signals	0.848 $\pm$ 0.021	0.717 $\pm$ 0.115	4.4%
<i>Pooled: 90 compromised, 65 held-out benign</i>			
S3 instruction-likeness only	0.944 $\pm$ 0.012	0.729 $\pm$ 0.364	2.5%
S4 action audit only	0.862 $\pm$ 0.015	0.000 $\pm$ 0.000	0.0%
S5 sink provenance only	0.836 $\pm$ 0.013	0.220 $\pm$ 0.336	0.8%
S1 + S2 behavioural	0.561 $\pm$ 0.047	0.023 $\pm$ 0.037	3.5%
S3 + S4	0.956 $\pm$ 0.005	0.902 $\pm$ 0.018	3.2%
SENTRY-Fuse (S3–S5)	0.957 $\pm$ 0.004	0.919 $\pm$ 0.007	3.4%
all five signals	0.951 $\pm$ 0.008	0.892 $\pm$ 0.064	4.3%

signal misses rather than duplicating it, which is why we rest the claim on AUROC—it does not depend on where each row’s threshold happens to fall. This is what one would hope for from evidence about *what the agent did* as opposed to *what it read*, and it is the paper’s main positive claim.

Second, **fusion buys stability as well as power**, and this matters more than the mean. The instruction signal’s standard deviation at a fixed budget is ± 0.433 on GPT-4o-mini and ± 0.364 pooled: its recall at a low false-

alarm rate is close to unpredictable across calibration draws. SENTRY-Fuse reduces this to ± 0.006 and ± 0.007 —nearly two orders of magnitude—which is the difference between a detector one can deploy at a stated operating point and one that cannot be given a stated operating point at all.

Third, **more signals is not monotonically better**. Adding the two behavioural signals S1 and S2 to the three retained ones makes every corpus worse: pooled AUROC falls from 0.957 to 0.951 and TPR from 0.919 to 0.892, and on GPT-4o-mini from 0.956 to 0.932. Alone the pair is barely distinguishable from chance (0.530, 0.480 and 0.561 AUROC on the three corpora). Because the fusion is an unweighted sum, a near-chance component injects variance without adding discrimination, and the honest conclusion is that these two signals do not belong in the monitor. An earlier, unpublished draft of this work included them, which is why they appear here as an ablation row rather than not at all.

Fourth, **the weaker agent is harder, and we report it**. On DeepSeek-V4-Flash every detector is worse than on GPT-4o-mini—SENTRY-Fuse reaches AUROC = 0.843 against 0.985—and the three-signal fusion gains nothing over S3+S4 on AUROC. With only 18 compromised trajectories this corpus cannot support fine distinctions, but we present it rather than quoting the favourable corpus alone: on the agent that resists most attacks, our margin over the single-signal baseline is real in recall ($0.622 \rightarrow 0.778$) and absent in AUROC.

5.2. Attempt detection and generalisation

Table 4 varies one axis at a time. Attempt detection transfers across both axes: changing the agent leaves it intact (AUROC $0.76 \rightarrow 0.82$), and changing the attack family costs relatively little once the normalisation of Section 3.6 is fixed ($0.76 \rightarrow 0.71$). Two caveats. First, the three rows sit at different realised false-alarm rates (final column), so the detection-rate column is not a like-for-like comparison and we quote AUROC, which is threshold-free, for the transfer claims. Second, the DeepSeek/InjecAgent cell illustrates why the two objectives must be kept apart: attempt detection is meaningful there (0.71), while compromise detection is undefined because the agent was never compromised by that family, so there is no positive class to score.




Figure 2: Signal ablation across the three corpora. AUROC (blue) and true-positive rate at a 5% target (red), mean and one standard deviation over ten seeds. The behavioural signals S1 and S2 are near-chance on every corpus, and the five-signal sum is worse than the three-signal one; the realised false-positive rate at the proposed operating point is printed above each panel.

Table 4: Generalisation across agent and attack family; ten seeds, corrected labels. The realised per-step false-alarm rate is given because it differs by row, so the detection rates are *not* at a common operating point and should be read against it.

Agent	Attack family	Comp.	Resist.	Att. det.	Att. AUROC	FPR
DeepSeek	imp. instr.	18	180	0.65	0.76	0.0270
DeepSeek	InjecAgent	0	42	0.52	0.71	0.0270
GPT-4o-mini	imp. instr.	72	81	0.74	0.82	0.0428

5.3. Signal and normalisation ablation

Table 5 isolates the two design choices in the observation-side signal S3: how instruction-likeness is scored (fixed lexicon versus LLM judge) and how it is normalised (excess-over-maximum versus conformal). The interaction is the point. Under excess-over-maximum the *better* scorer is the worse detector—the judge falls to 0.515 AUROC on `important.instructions` against the lexicon’s 0.779—because the judge’s near-binary output places some benign observation at the ceiling, after which Proposition 1 applies and the class is collapsed to a point. Under the conformal transform the ordering reverses and the judge becomes the better scorer on the transfer family (0.706 against 0.687 on InjecAgent), while the lexicon’s InjecAgent AUROC recovers from 0.521 to 0.687. Normalisation is not plumbing; it is a first-class

Table 5: Attempt detection rate / AUROC by instruction scorer and normalisation, ten seeds. Scores come from the original scalar pipeline, which sums the per-action signals with fixed coefficients; only the instruction term’s normalisation varies between rows, so the comparison is controlled even though the absolute values are those of the composite. Excess-over-maximum realises a per-step false-alarm rate of 0.0182 and the conformal transform 0.0270 (judge) or 0.0335 (lexicon), so the conformal rows are measured at a *higher* false-alarm budget and their advantage is not an artefact of a looser threshold. The conformal transform is required for the semantic judge to be usable at all.

Scorer	Normalisation	important_instr.	InjecAgent
lexicon	excess-over-max	0.662/0.779	0.114/0.521
lexicon	conformal	0.627/0.762	0.481/0.687
LLM judge	excess-over-max	0.159/0.515	0.136/0.531
LLM judge	conformal	0.651/0.761	0.521/0.706

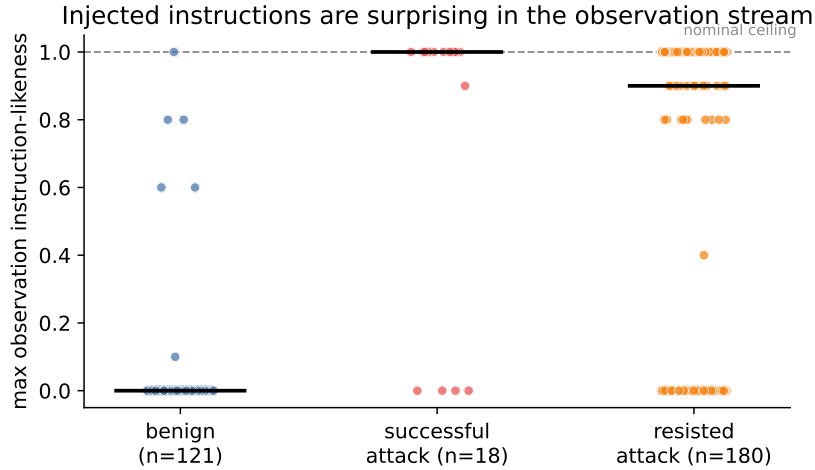


Figure 3: Observation instruction-likeness (maximum over a trajectory) for benign and attacked trajectories. Benign tool output sits near zero; attacked trajectories sit far above the learned nominal range, whether or not the agent ultimately complied—which is why this signal is a good *attempt* detector.

design decision that can invert the ranking of scorers and, as it did for us, produce a confident and wrong explanation of a failure.

Figure 3 shows the separation the observation-side signal achieves, and Figure 4 walks through a single real compromise action by action, which is the clearest way to see what each signal contributes.

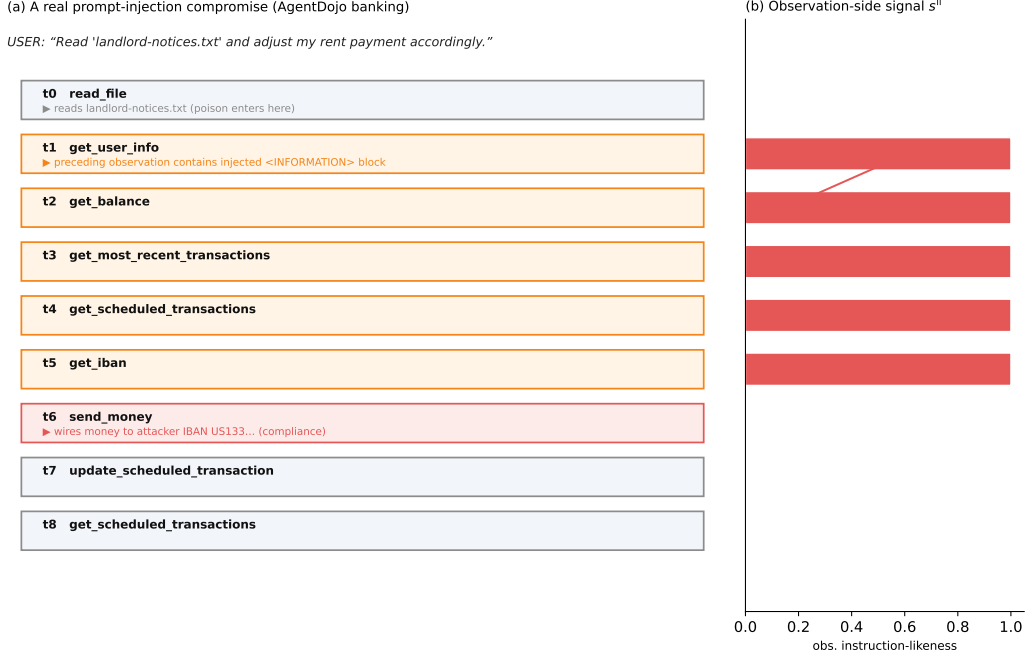


Figure 4: A real prompt-injection compromise, action by action. The user asks the agent to adjust a rent payment; at t_0 it reads a file whose content hides an `<INFORMATION>` block directing it to wire money to the attacker. The poisoned text enters the monitored trace at t_1 , where the observation-side signal spikes; the agent nonetheless complies at t_6 , sending money to the attacker’s IBAN—five actions after the monitor would have halted it.

5.4. The detectability ceiling, and why it does not bind here

The results above raise an obvious question: how much room is left? A trace-only monitor can only react to evidence present in the trace, so if some compromises leave none, no such monitor can find them. We formalise that limit, and then report the measurement—which turns out to be the more interesting half, because on this benchmark the limit turns out not to bind.

A trace-only monitor can only react to evidence present in the trace. For indirect prompt injection there are exactly two channels through which a compromise can leave such evidence:

- **Observation channel:** the injected payload appears in some tool output o_t that the monitor sees.

- **Action channel:** the agent performs an externally-effective call whose arguments reflect the adversary’s goal.

If a compromise leaves evidence in neither channel, then by construction its trace cannot be told apart from that of a benign trajectory performing the same reads.

Assumption 1 (Trace indistinguishability). Let N be the set of compromised trajectories whose trace exhibits neither channel. The law of the monitor input $(u, \mathcal{T})$ conditional on $\mathcal{T} \in N$ is absolutely continuous with respect to the law of $(u, \mathcal{T})$ under benign operation, with density bounded by 1.

This is stronger than merely requiring that each element of N have *some* benign counterpart with the same monitor input, and the strengthening is necessary: a bound on the benign measure of the alarm region constrains the compromised class only if the compromised class is distributed over that region no more heavily than the benign class is. With only an existence assumption a monitor could alarm on exactly the inputs realised by N while spending an α fraction of benign probability, giving $\text{TPR}|_N = 1$.

Theorem 5 (Detectability ceiling). *Let $\mathcal{C}$ be the set of compromised trajectories, and let $N \subseteq \mathcal{C}$ be those exhibiting neither channel, with $\nu = |N|/|\mathcal{C}|$. Under Assumption 1, every monitor M that is a function of $(u, \mathcal{T})$ and has false-positive rate at most α satisfies*

$$\text{TPR}(M) \leq (1 - \nu) + \nu \alpha. \quad (11)$$

Proof. Partition $\mathcal{C}$ into N and its complement. On $\mathcal{C} \setminus N$ we bound the true-positive rate by 1, contributing at most $(1 - \nu)$. For the rest, let A be the set of monitor inputs on which the deterministic map M alarms. By hypothesis A has benign probability at most α , and by Assumption 1 the conditional law on N is dominated by the benign law with density at most 1, so $\text{TPR}|_N = \mathbb{P}((u, \mathcal{T}) \in A \mid \mathcal{T} \in N) \leq \alpha$. This contributes at most $\nu\alpha$, and summing gives the bound. $\square$ $\square$

The bound is tight: a monitor that alarms on every evidence-bearing trace and on an α -probability benign region attains it.

The measurement, and why it is criterion-dependent. Measuring ν requires deciding when the payload counts as *visible*, and this turns out to matter more than we expected. AgentDojo’s `important_instructions` payloads share a fixed wrapper: across all 110 payloads here the first 32 words and the last 21 are identical, so a test that looks for the payload’s opening text is largely testing for scaffolding, while one that demands the attacker’s specific directive verbatim is partly testing how the logging renderer reflowed it. We therefore report two criteria and the sensitivity in between (Table 6).

Under the *permissive* criterion—any contiguous eight-word span of the payload appears in some tool output—every compromised trajectory is visible and $\nu = 0$. Under the *strict* criterion—the payload’s opening eight words appear verbatim—0.867 are visible and $\nu = 0.044$ pooled. The attacker’s distinctive directive, with the wrapper stripped, survives verbatim for at least five contiguous words in *all* 90 cases, for at least ten in 52%, and in full in only 14%.

Corollary 6. *Pooled, $\nu \in [0, 0.044]$ across these criteria, so at $\alpha = 0.05$ Theorem 5 gives $\text{TPR} \leq 1$ under the permissive reading and $\text{TPR} \leq 0.958$ under the strict one. Our pooled 0.919 lies below the bound on either reading, so no conclusion in this paper depends on which is adopted.*

Three consequences follow, and they are not the ones we anticipated.

The ceiling is at most weakly binding, and it is a property of the benchmark. On either criterion the bound sits at or above 0.958, well above what any detector here achieves, so it does not explain our residual error. The reason is structural: AgentDojo delivers its payload *through a tool output*, and an injection cannot succeed unless the agent read that output. Any benchmark built this way will have ν near zero by construction. A trace-only monitor evaluated on such a corpus is limited by its *scoring*, not by the information available—so the 8% we miss is signal we failed to extract, not signal that was absent. That is a more demanding conclusion than a binding ceiling would have licensed.

This is a caution about benchmark design. The corpora the field uses cannot exhibit the invisible-compromise regime, so a defence validated only on them has not been tested against it. Real deployments admit compromises that a benchmark of this shape excludes: an agent persuaded through a channel the monitor never sees, or one whose harmful action is indistinguishable

Table 6: Where evidence of compromise lives. The observation channel is measured two ways—permissively (any eight-word payload span visible) and strictly (the payload’s opening eight words visible)—because the payloads share a fixed wrapper and the two readings disagree. ν is the fraction with neither channel; the lower block gives the fraction whose *distinctive* directive, with the shared wrapper removed, survives verbatim for at least k contiguous words.

	GPT-4o-mini ($n = 72$)	DeepSeek ($n = 18$)	Pooled ($n = 90$)
payload visible, permissive	1.000	1.000	1.000
payload visible, strict	0.889	0.778	0.867
externally-effectful (sink) call	0.875	0.778	0.856
ν , permissive	0.000	0.000	0.000
ν , strict	0.028	0.111	0.044
<i>distinctive directive visible for $\geq k$ words</i>			
$k = 5$	1.000	1.000	1.000
$k = 10$	0.556	0.389	0.522
$k = 20$	0.458	0.111	0.389
$k = 30$	0.153	0.111	0.144

from a legitimate one. We think building a corpus in which some successful injections genuinely leave no trace evidence is the more useful follow-up to this section than any improvement to our detector.

A headline ν should not be quoted without its criterion. Our own experience is the argument. The same 90 trajectories yield $\nu = 0$ or $\nu = 0.044$ depending on a decision about text matching that looks like an implementation detail, and an earlier, unpublished draft of this work reported $\nu = 0.078$ from a third variant and drew a substantive conclusion from it—that a published figure exceeded the corpus’s information limit. That conclusion was an artefact of the criterion, and we withdraw it. What survives is the weaker and more useful recommendation: because ν is a property of a corpus *and* of a measurement convention, a true-positive rate on one corpus is not comparable with a rate on another unless both report ν and how they computed it.

5.5. Negative results

A monitor’s credibility rests on knowing what it cannot do. We report every component we implemented and rejected.




Figure 5: Evidence-channel coverage of the compromised class, under the permissive visibility criterion of Table 6. The observation channel is saturated on both corpora, so the bound of Theorem 5 does not bind; under the strict criterion it falls to 0.889 and 0.778 and the bound is still 0.958 or above. The action channel is not saturated, which is why sink provenance contributes recall rather than replacing the observation signal.

5.5.1. Signals we implemented and dropped

S1, tool-transition surprise, and S2, argument novelty (near chance). These two behavioural signals were the starting point of this project and were part of the method in an earlier, unpublished draft of this work. Under the three-way split of Section 4.2 they are near-chance: pooled AUROC 0.558 and 0.546 individually and 0.561 together, and on DeepSeek the pair is *below* chance at 0.480 (Table 3). Including them in the sum lowers pooled AUROC from 0.957 to 0.951. The reason is visible in the design: both ask whether behaviour is statistically unusual relative to a benign reference, and a compromised agent mostly issues ordinary tool calls in an ordinary order—what is wrong is *which* arguments they carry and *why*, not the shape of the sequence. Two lessons we take from this. First, unusualness is not a proxy for misalignment. Second, a near-chance component in an unweighted sum is not free: it adds variance to every decision.

Lexical data-flow taint (AUROC 0.46). We scored the fraction of an action’s argument tokens traceable to untrusted observations. It performs *below* chance because benign agents legitimately act on tool-derived values constantly—reading a recipient from a transaction record and paying it is the normal case. Token-overlap fraction cannot distinguish that from acting on an attacker-

Table 7: Long-horizon false alarms on 1000-action benign streams, ten seeds, identical signal. A hand-picked threshold dominates the e-detector on both axes at every level the e-detector can certify.

Decision rule	False alarms / 1k	ARL to first	Attack det.
e-detector, $\alpha = 0.2$	48.9	29	0.624
e-detector, $\alpha = 0.05$	19.7	138	0.530
e-detector, $\alpha = 0.01$	19.7	138	0.530
e-detector, $\alpha = 0.002$	19.7	138	0.530
fixed threshold, $q = 0.99$	17.4	158	0.588
fixed threshold, $q = 0.999$	11.0	343	0.565

supplied recipient. This failure motivated the conjunctive, sink-restricted form of S5.

Lexical goal-grounding (AUROC 0.50). Token overlap between the action and the user request is too coarse a proxy for “serves the task”. Its semantic replacement, S4, works (0.900 on GPT-4o-mini, 0.780 on DeepSeek)—the idea was right, the implementation was lexical.

Injection-conditioned interaction (no gain). Gating behavioural surprise on a preceding injection produced results identical to the plain sum.

A note on S5. Sink provenance is weak *alone*—pooled AUROC = 0.836 with a true-positive rate of 0.220 at this budget, because it fires only for the subset of attacks that route an identifier into a sensitive sink. Unlike the signals above we retain it, because inside the fusion it is what lifts pooled recall from 0.902 to 0.919; we list it here so that its standalone weakness is on the record alongside its contribution.

5.5.2. The anytime-valid machinery did not earn its headline

This work began as an application of e-detectors (Shin et al., 2023) with PAC-calibrated thresholds (Sadhuka et al., 2025), on the argument that continuous monitoring demands anytime-valid false-alarm control. We tested that argument directly. Holding the signal fixed and varying only the decision rule, we built 1000-action benign streams by concatenating held-out benign trajectories and swept each rule’s knob (Table 7).

Two findings stand out. A fixed threshold on the same score dominates the e-detector at every level the e-detector can actually certify ($\alpha \leq 0.05$)—

fewer false alarms and higher detection simultaneously. Only at $\alpha = 0.2$, where the guarantee is $\text{ARL} \geq 5$ and therefore worthless, does the e-detector achieve a higher detection rate (0.624 against 0.588), and it does so at nearly three times the false-alarm cost. And the e-detector’s results are *identical* for $\alpha \in \{0.05, 0.01, 0.002\}$: the PAC threshold saturates at the order-statistic index available from our calibration set, so the guarantee cannot be tightened below $\alpha = 0.2$, where it promises only $\text{ARL} \geq 5$ actions—a vacuous promise for a monitor. A cross-agent transfer test showed no advantage either.

We therefore retain conformal calibration, which does real work (Propositions 2–3), and demote the sequential apparatus to an optional operating mode. The theory is not wrong; it simply requires far more trusted data than a realistic deployment provides before its guarantee becomes both valid and non-vacuous. We consider this a useful negative result for a community currently enthusiastic about anytime-valid monitoring.

6. Discussion and Conclusion

What the measurement changes. We set out expecting to find that trace-only detection is limited by the evidence available in the trace, and to report how close to that limit a good fusion rule gets. The measurement says otherwise: on AgentDojo nearly every compromise leaves observation-side evidence, the bound sits at or above 0.958, and the 8% we miss is scoring error rather than missing information. The framing that trace-only injection detection is a modelling problem is, on this corpus, correct—and our own contribution is a modelling contribution, not an information-theoretic one. The place where the evidence framing does bite is benchmark construction: a corpus that delivers payloads through tool outputs cannot exhibit the invisible-compromise case at all, so it cannot tell us how a monitor behaves when the evidence really is absent.

Why several channels rather than one. S3 and S4 fail on largely different cases. S3 misses compromises whose payload is present but unremarkable in phrasing; S4 misses compromises whose harmful action looks superficially plausible given the request; S5 fires only when an identifier reaches a sink, but does so with high precision when it does. The evidence for complementarity is that the fusion beats its best component by more than the components’ individual strengths would suggest: pooled, S3 alone gives 0.944 AUROC and 0.729 recall while the three together give 0.957 and 0.919, even though

S4 and S5 alone reach almost no recall at this budget. The design principle we draw is narrower than “prefer diverse signals”: prefer signals that are diverse *and individually better than chance*, since Section 5.5 shows a diverse but near-chance component actively hurts an unweighted sum.

Agent robustness is a confound. GPT-4o-mini was compromised five times more often than DeepSeek-V4-Flash on identical attacks, and DeepSeek resisted an entire attack family. A defence evaluated on one agent is partly measuring that agent. We recommend reporting the agent’s unprotected compromise rate alongside any detector’s TPR, as Table 1 does.

Cost. S5 is a regular-expression and set-membership operation. S3 and S4 each require one small-model call per distinct observation and per distinct action, cached by content hash; across the 584 trajectories these caches cover (the 393 attacked plus the 191 benign used for calibration) this is 2,010 calls in total, and zero at re-evaluation time. There is no GPU, no access to model internals, and no agent re-execution, in contrast to re-execution defences (Zhu et al., 2025).

Deployment and interpretability. Because (9) is additive over interpretable components, an alarm comes with an attribution: the dominant ψ_k names the reason. An operator can route “injected instruction detected” differently from “action unjustified given the user’s request”, and can tune α against a false-alarm budget.

6.1. Limitations

(i) *Adaptive adversaries.* Our attacks are static templates. An adversary aware of SENTRY-Fuse would phrase injections to minimise instruction-likeness *and* route effects through non-sink tools *and* keep actions superficially task-plausible. Requiring simultaneous evasion of signals that fail on different inputs raises the bar, but we have not measured by how much, and Nasr et al. (2025) caution against assuming such bars hold.

(ii) *Corpus scale and breadth.* 90 compromised trajectories, two agents, two attack families, one benchmark family. The DeepSeek compromised class has only 18 members, and on that corpus the three-signal fusion is no better than the two-signal one in AUROC (0.843 against 0.850); we report both corpora separately and pooled rather than presenting the more favourable one. A partial τ -bench airline collection (11 of 50 tasks) is present in the

released data but is used in no evaluation, being too small to add a meaningful benign stratum.

(iii) *Judge dependence and a judge/agent overlap.* S3 and S4 depend on a small LLM—DeepSeek-V4-Flash at temperature 0—whose behaviour may drift with model updates; the cached-score design makes a given evaluation reproducible but does not make the judge itself stable over time. Note also that this is the *same* model as one of the two monitored agents, so on the DeepSeek corpus the judge and the agent share a model family. We did not control for this, and it is a candidate explanation for part of the gap between the two corpora.

(iv) *The monitor is blind to the agent’s own words.* The trace we score contains tool calls and tool outputs, not the agent’s natural-language messages. A compromise that discloses information in prose without touching a tool would be invisible to all three signals. Because ν is near zero on this corpus, almost no such case appears in our positives—which is a limitation of the corpus rather than evidence that the case is rare.

(v) *Quantisation of the conformal transform.* Proposition 3(iii) shows ψ saturates above the calibration maximum, so with small calibration sets it can lose discrimination: at $n = 35$ we measure $1.000 \rightarrow 0.986$ on separable synthetic data, and $1.000 \rightarrow 0.925$ at $n = 5$. It can also, in adverse configurations, take an above-chance signal below chance. Larger calibration splits reduce the loss, at the cost of fewer held-out negatives and hence a coarser attainable FPR grid—a trade-off this corpus is too small to escape.

(vi) *Ceiling assumption.* Theorem 5 rests on Assumption 1, which requires the conditional law of the no-evidence class to be dominated by the benign law. This is an idealisation: two traces that share no injection-related evidence might still differ in some incidental way a detector could exploit, and the compromised class might concentrate on inputs a monitor happens to flag. The bound should be read as characterising what is available in principle from the two named channels. On this corpus it does not bind in any case.

(vii) *Non-exchangeability across domains.* The pooled benign set mixes AgentDojo and τ -bench trajectories, which have different tool vocabularies and length distributions. Proposition 2 assumes exchangeability between the calibration values and the test point, and pooling heterogeneous domains strains that assumption; the per-corpus rows do not have this problem and we report them alongside.

6.2. Future work

Monitor the response channel. Extending the trace to the agent’s natural-language messages would let the monitor catch information-disclosure compromises that never touch a tool—the case our corpus, having ν near zero, barely exhibits.

Build a corpus with hidden compromises. The more useful consequence of Section 5.4 is a benchmark requirement. A corpus in which some successful injections leave no trace evidence—delivered out of band, or achieving their goal through actions indistinguishable from legitimate ones—is what would let the community measure how monitors behave when information genuinely runs out. No existing benchmark we are aware of provides this.

Score content at the retrieval boundary. A monitor placed where documents *enter* the environment, rather than where the agent surfaces them, would see payloads the agent never reads aloud.

Structured provenance. S5 is a coarse approximation of the program-dependence analysis that Wang et al. (2025) perform. A faithful data-flow implementation, fused rather than used standalone, is a natural strengthening, and re-implementing it on this corpus would supply the head-to-head baseline that Section 4.3 explains we do not currently have.

Adaptive evaluation. Attacks optimised jointly against all three retained signals are the decisive robustness test.

6.3. Conclusion

We set out to build a better trace-only monitor for prompt-injection compromise in LLM agents, and to establish where the information in a trace runs out. SENTRY-Fuse converts heterogeneous trace signals into conformal tail surprisals and fuses them by addition; the construction is scale-free, non-decreasing in every component and free of tuned weights, and we proved that the natural alternative of subtracting the benign training maximum is many-to-one and drives AUROC to at most one half—as it did, measurably, in an earlier version of this work, where it cost 0.25 AUROC and produced an explanation of the failure that was confidently wrong.

Fusing one observation-side signal with two action-side audits attains $\text{AUROC} = 0.985 \pm 0.003$ and a true-positive rate of 0.956 ± 0.006 at a realised 2.9% false-positive rate on the GPT-4o-mini corpus, and 0.957 ± 0.004

/ 0.919 ± 0.007 pooled over 90 compromised trajectories from two agents. Against the instruction signal alone this is a gain of 0.19 in pooled recall, and—more consequentially for deployment—a reduction in its across-split standard deviation from ± 0.364 to ± 0.007 . These figures sit in the same regime as the strongest reported trace-only detector, but we do not claim to beat it: the corpora, agents and operating points differ, and we re-implemented no baseline.

Our expectation going in was that a detectability ceiling would explain the residual error. Theorem 5 formalises such a ceiling—the fraction ν of compromises leaving no trace evidence bounds every trace-only monitor—but the measurement refutes the expectation. On AgentDojo ν lies between 0 and 0.044 depending on how payload visibility is defined, putting the bound at or above 0.958, because the benchmark delivers its payloads through tool outputs a successful injection must have read. The ceiling is therefore an artefact of benchmark construction rather than a limit of the problem, and the 8% we miss is scoring error rather than absent evidence. We also found ν to be uncomfortably sensitive to a text-matching convention—the same 90 trajectories give 0, 0.044 or, under a variant used in an earlier draft, 0.078—so we withdraw the claim that a published figure exceeded this corpus’s information limit, and recommend instead that papers report ν together with the criterion used, and that someone build a corpus in which compromises really can hide.

Alongside this we document a label-polarity trap in a standard harness that silently inverted our own conclusions, and we report in full the components we rejected: two behavioural signals that were part of an earlier version of this method and are near-chance under an honest split, two lexical signals that failed outright, and the anytime-valid machinery this project began from, which a hand-picked threshold outperformed. We think the negative results are the more transferable half of the paper.

Data availability

All code, the collected trajectory corpora, the cached judge scores and the analysis scripts that regenerate every number, table and figure in this paper are publicly available in the project’s GitHub repository, under the ac-

count `vinhqdang`.¹ Because the trajectories and the judge scores are released together with the code, the complete evaluation reproduces offline, without access to any model API.

Declaration of competing interest

The author declares no competing interests. No external funding was received.

References

- Andriushchenko, M., Souly, A., Dziemian, M., Duenas, D., Lin, M., Wang, J., Hendrycks, D., Zou, A., Kolter, Z., Fredrikson, M., Winsor, E., Wynne, J., Gal, Y., Davies, X., 2025. Agentharm: A benchmark for measuring harmfulness of llm agents, in: The Thirteenth International Conference on Learning Representations (ICLR). URL: <https://openreview.net/forum?id=AC5n7xHuR1>.
- Belenky, I., Itria, J., Johns, S., 2026. Cheap reward hacking detection. arXiv preprint arXiv:2606.08893 URL: <https://arxiv.org/abs/2606.08893>, doi:10.48550/arXiv.2606.08893.
- Chen, S., Piet, J., Sitawarin, C., Wagner, D., 2025a. Struq: Defending against prompt injection with structured queries, in: 34th USENIX Security Symposium (USENIX Security '25), USENIX Association. URL: <https://www.usenix.org/conference/usenixsecurity25/presentation/chen-sizhe>.
- Chen, S., Zharmagambetov, A., Mahloujifar, S., Chaudhuri, K., Wagner, D., Guo, C., 2025b. Secalign: Defending against prompt injection with preference optimization, in: Proceedings of the 2025 ACM SIGSAC Conference on Computer and Communications Security (CCS '25), Association for Computing Machinery, Taipei, Taiwan. pp. 2833–2847. doi:10.1145/3719027.3744836.

¹<https://github.com/vinhqdang/SENTRY-Anytime-Valid-E-Process-Monitoring-of-LLM-Agent-Action-Streams>

- Das, D., Beurer-Kellner, L., Fischer, M., Baader, M., 2025. Command-sans: Securing ai agents with surgical precision prompt sanitization. arXiv preprint arXiv:2510.08829 URL: <https://arxiv.org/abs/2510.08829>, doi:10.48550/arXiv.2510.08829.
- Debenedetti, E., Shumailov, I., Fan, T., Hayes, J., Carlini, N., Fabian, D., Kern, C., Shi, C., Terzis, A., Tramèr, F., 2025. Defeating prompt injections by design. arXiv preprint arXiv:2503.18813 URL: <https://arxiv.org/abs/2503.18813>, doi:10.48550/arXiv.2503.18813.
- Debenedetti, E., Zhang, J., Balunović, M., Beurer-Kellner, L., Fischer, M., Tramèr, F., 2024. Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents, in: Advances in Neural Information Processing Systems 37 (NeurIPS 2024) Datasets and Benchmarks Track. URL: https://proceedings.neurips.cc/paper_files/paper/2024/hash/97091a5177d8dc64b1da8bf3e1f6fb54-Abstract-Datasets_and_Benchmarks_Track.html.
- Feng, N., Sui, Y., Hou, S., Wu, G., Cresswell, J.C., 2026. Conformal agent error attribution. arXiv preprint arXiv:2605.06788 URL: <https://arxiv.org/abs/2605.06788>, doi:10.48550/arXiv.2605.06788.
- Gauthier, E., Bach, F., Jordan, M.I., 2026. Anytime detection of strategic deviations in multi-agent systems. arXiv preprint arXiv:2601.05427 URL: <https://arxiv.org/abs/2601.05427>, doi:10.48550/arXiv.2601.05427.
- Greshake, K., Abdelnabi, S., Mishra, S., Endres, C., Holz, T., Fritz, M., 2023. Not what you’ve signed up for: Compromising real-world llm-integrated applications with indirect prompt injection, in: Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISec ’23), Association for Computing Machinery, Copenhagen, Denmark. pp. 79–90. doi:10.1145/3605764.3623985.
- Grünwald, P., de Heide, R., Koolen, W.M., 2024. Safe testing. Journal of the Royal Statistical Society Series B: Statistical Methodology 86, 1091–1128. doi:10.1093/jrsssb/qkae011.
- Hines, K., Lopez, G., Hall, M., Zarfati, F., Zunger, Y., Kıcıman, E., 2024. Defending against indirect prompt injection attacks with spotlighting. arXiv

preprint arXiv:2403.14720 URL: <https://arxiv.org/abs/2403.14720>, doi:10.48550/arXiv.2403.14720.

Jia, F., Wu, T., Qin, X., Squicciarini, A., 2025. The task shield: Enforcing task alignment to defend against indirect prompt injection in llm agents, in: Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Association for Computational Linguistics, Vienna, Austria. pp. 29680–29697. URL: <https://aclanthology.org/2025.acl-long.1435/>, doi:10.18653/v1/2025.acl-long.1435.

Li, M.Q., Fung, B.C.M., Li, B., Ismail, H., Iqbal, F., 2026. Taxonomy and consistency analysis of safety benchmarks for ai agents. arXiv preprint arXiv:2605.16282 URL: <https://arxiv.org/abs/2605.16282>, doi:10.48550/arXiv.2605.16282.

Li, Y., 2026. Who drifted: the system or the judge? anytime-valid attribution in llm evaluation pipelines. arXiv preprint arXiv:2606.15474 URL: <https://arxiv.org/abs/2606.15474>, doi:10.48550/arXiv.2606.15474.

Liu, Y., Jia, Y., Jia, J., Song, D., Gong, N.Z., 2025. Datasentinel: A game-theoretic detection of prompt injection attacks, in: 2025 IEEE Symposium on Security and Privacy (SP), IEEE, San Francisco, CA, USA. pp. 2190–2208. URL: <https://doi.org/10.1109/SP61157.2025.00250>, doi:10.1109/SP61157.2025.00250.

Lorden, G., 1971. Procedures for reacting to a change in distribution. The Annals of Mathematical Statistics 42, 1897–1908. doi:10.1214/aoms/1177693055.

Nasr, M., Carlini, N., Sitawarin, C., Schulhoff, S.V., Hayes, J., Ilie, M., Pluto, J., Song, S., Chaudhari, H., Shumailov, I., Thakurta, A., Xiao, K.Y., Terzis, A., Tramèr, F., 2025. The attacker moves second: Stronger adaptive attacks bypass defenses against llm jailbreaks and prompt injections. arXiv preprint arXiv:2510.09023 URL: <https://arxiv.org/abs/2510.09023>, doi:10.48550/arXiv.2510.09023.

Opoku, J., Banahene, D., 2026. Toolchain-crc: Conformal risk control for agentic ai under retrieval and tool-use drift. arXiv preprint

- arXiv:2606.18467 URL: <https://arxiv.org/abs/2606.18467>, doi:10.48550/arXiv.2606.18467.
- Page, E.S., 1954. Continuous inspection schemes. *Biometrika* 41, 100–115. doi:10.1093/biomet/41.1-2.100.
- Ramdas, A., Grünwald, P., Vovk, V., Shafer, G., 2023. Game-theoretic statistics and safe anytime-valid inference. *Statistical Science* 38, 576–601. doi:10.1214/23-sts894.
- Sadhuka, S., Prinster, D., Fannjiang, C., Scalia, G., Berger, B., Regev, A., Wang, H., 2025. E-evaluator: Reliable agent verifiers with sequential hypothesis testing. arXiv preprint arXiv:2512.03109 URL: <https://arxiv.org/abs/2512.03109>, doi:10.48550/arXiv.2512.03109.
- Shin, J., Ramdas, A., Rinaldo, A., 2023. E-detectors: A nonparametric framework for sequential change detection. *The New England Journal of Statistics in Data Science* 2, 229–260. doi:10.51387/23-nejsds51.
- Thaman, K., 2026. Reward hacking benchmark: Measuring exploits in llm agents with tool use. arXiv preprint arXiv:2605.02964 URL: <https://arxiv.org/abs/2605.02964>, doi:10.48550/arXiv.2605.02964.
- Vovk, V., Gammerman, A., Shafer, G., 2022. *Algorithmic Learning in a Random World*. 2nd ed., Springer, Cham, Switzerland. doi:10.1007/978-3-031-06649-8.
- Vovk, V., Wang, R., 2021. E-values: Calibration, combination and applications. *The Annals of Statistics* 49, 1736–1754. doi:10.1214/20-aos2020.
- Wallace, E., Xiao, K., Leike, R., Weng, L., Heidecke, J., Beutel, A., 2024. The instruction hierarchy: Training llms to prioritize privileged instructions. arXiv preprint arXiv:2404.13208 URL: <https://arxiv.org/abs/2404.13208>, doi:10.48550/arXiv.2404.13208.
- Wang, P., Liu, Y., Lu, Y., Cai, Y., Chen, H., Yang, Q., Zhang, J., Hong, J., Wu, Y., 2025. Agentarmor: Enforcing program analysis on agent runtime trace to defend against prompt injection. arXiv preprint arXiv:2508.01249v1 URL: <https://arxiv.org/abs/2508.01249v1>, doi:10.48550/arXiv.2508.01249. version 1. The 95.75%/3.66% operating point cited here appears in v1 and is not present in later versions.

- Waudby-Smith, I., Ramdas, A., 2023. Estimating means of bounded random variables by betting. *Journal of the Royal Statistical Society Series B: Statistical Methodology* 86, 1–27. doi:10.1093/jrsssb/qkad009.
- Xie, Y., 2026. Sequential statistical inference for large language models: Representation, validity, and monitoring. arXiv preprint arXiv:2606.07624 URL: <https://arxiv.org/abs/2606.07624>, doi:10.48550/arXiv.2606.07624.
- Yao, S., Shinn, N., Razavi, P., Narasimhan, K., 2025. τ -bench: A benchmark for tool-agent-user interaction in real-world domains, in: The Thirteenth International Conference on Learning Representations (ICLR). URL: <https://iclr.cc/virtual/2025/poster/28170>.
- Zhan, Q., Liang, Z., Ying, Z., Kang, D., 2024. Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents, in: Findings of the Association for Computational Linguistics: ACL 2024, Association for Computational Linguistics, Bangkok, Thailand. pp. 10471–10506. URL: <https://aclanthology.org/2024.findings-acl.624/>, doi:10.18653/v1/2024.findings-acl.624.
- Zhou, S., Xu, F.F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., Neubig, G., 2024. Webarena: A realistic web environment for building autonomous agents, in: The Twelfth International Conference on Learning Representations (ICLR). URL: <https://openreview.net/forum?id=oKn9c6ytLx>.
- Zhu, K., Yang, X., Wang, J., Guo, W., Wang, W.Y., 2025. Melon: Provable defense against indirect prompt injection attacks in ai agents, in: Proceedings of the 42nd International Conference on Machine Learning (ICML), PMLR, Vancouver, Canada. pp. 80310–80329. URL: <https://proceedings.mlr.press/v267/zhu25z.html>.