

Breaking the Observation Ceiling: Conformal Fusion of Trace Signals for Detecting Prompt-Injection Compromise in LLM Agents

Quang-Vinh Dang^a

^a*British University Vietnam, Hung Yen, Vietnam*

Abstract

Large language model (LLM) agents that call external tools ingest untrusted content and fold it into the context that drives their next action, exposing them to indirect prompt injection. Runtime monitors that watch the agent’s black-box execution trace are attractive because they are model-agnostic and cheap, but the literature has converged on a single family of evidence: detecting whether untrusted text *looks like an instruction*. We show that this choice imposes a hard information limit and that crossing it requires a different kind of evidence.

We make four contributions. First, we introduce **SENTRY-Fuse**, a monitor that converts heterogeneous trace signals—two that score the *observation* stream and two that score the *action* stream—into conformal tail surprisals against a benign calibration split and fuses them additively. The construction is scale-free, strictly monotone in every component, and requires no tuned weights; we prove that these properties are exactly what a fusion rule needs, and we exhibit a widely-used alternative normalisation that provably destroys signal. Second, we show that action-side evidence is *complementary* rather than redundant: on the pooled corpus SENTRY-Fuse attains a true-positive rate of 0.902 at a 3.66% false-positive rate, strictly above the 0.867 attainable from payload visibility alone, and on a matched GPT-4o-mini corpus it reaches 0.960 ± 0.010 , exceeding the strongest published trace-only figure (0.9575). Third, we characterise the **detectability ceiling** of trace-only monitoring: over 90 genuinely compromised trajectories, 7.8% leave no trace evidence whatsoever, bounding every such monitor at 0.922—a bound that

Email address: `vinh.dq4@buv.edu.vn` (Quang-Vinh Dang)

published numbers above 0.92 on this benchmark exceed and must therefore explain. Fourth, we document a label-polarity pitfall in the standard AgentDojo harness that inverts the compromised and resisted classes, which silently inverted our own results, and we give the ground-truth procedure that settles it. We report our negative results in full, including three rejected signals and the finding that a hand-tuned threshold dominates the anytime-valid e-detector we began from. All code, trajectories and scripts are released.

Keywords: LLM agents, indirect prompt injection, runtime monitoring, conformal prediction, anomaly detection, AI agent security

1. Introduction

Autonomous agents built on large language models now read email, browse the web, query databases and move money on a user’s behalf. To act, they consume *untrusted* content—web pages, documents, tool outputs—and fold it directly into the context that determines their next tool call. This is the attack surface exploited by *indirect prompt injection*: an adversary plants an instruction inside data the agent will later read, and the agent, unable to separate trusted instructions from untrusted data, follows it (Zhan et al., 2024; Debenedetti et al., 2024). As agents move from demos into production, a runtime monitor—something that watches the live execution trace and halts or escalates when behaviour departs from what is safe—becomes as important as the agent itself.

Monitors that observe only the black-box trace are especially attractive in practice. They need no access to model weights or activations, they impose no retraining cost, they compose with whichever model a deployment happens to use, and they can be dropped in front of a third-party agent. A substantial line of work therefore builds detectors over traces (Wang et al., 2025; Liu et al., 2025; Jia et al., 2025). Almost all of it, however, rests on one kind of evidence: examining the untrusted content the agent read and asking whether it *looks like an instruction*. Keyword and perplexity heuristics do this lexically; fine-tuned detectors and LLM judges do it semantically.

The observation that started this work.. We began from that same premise and built a monitor on it. It performed well on the standard benchmark attack and then collapsed, to chance, when we changed the attack template. Our first explanation was that the new attack was phrased too subtly. That

explanation was *wrong*: the new payloads were, if anything, more imperative, and they appeared in the trace more often. The collapse came from our own normalisation step. Diagnosing that failure honestly led us to a more general question, which organises this paper:

What, exactly, can a monitor learn from a black-box agent trace, and where does that information run out?

Answering it requires separating two objectives that the literature routinely conflates. *Attempt detection* asks whether untrusted instructions entered the agent’s context at all; *compromise detection* asks whether the agent actually did the attacker’s bidding. They have different positive classes, different ceilings, and—as we show—different optimal evidence. A monitor tuned for the first is not automatically good at the second.

Contributions.. Our contributions are as follows.

1. **SENTRY-Fuse, a conformal fusion monitor** (Section 4). We combine four trace signals—two scoring the observation stream and two scoring the action stream—by mapping each to its conformal tail surprisal against a benign calibration split and summing. We show (Propositions 2–3) that the construction is distribution-free valid, scale-free and strictly monotone in every component, and that the excess-over-maximum normalisation used in an earlier version of this work is many-to-one and can annihilate an entire attack family—which it did, measurably (Section 4.2).
2. **Action-side evidence is complementary, not redundant** (Section 7). Signals that audit *what the agent did* carry information that signals scoring *what the agent read* cannot. The fused monitor exceeds the payload-visibility ceiling, and on a matched corpus it exceeds the strongest published trace-only number.
3. **The detectability ceiling of trace-only monitoring** (Section 8). We formalise and measure the limit: some successful injections leave no trace evidence at all, capping any trace-only detector. Theorem 5 makes the bound precise; empirically it is 0.922, and we reach 97.8% of it.
4. **Methodological findings and honest negatives** (Sections 6 and 9). We document an label-polarity trap in the standard evaluation harness

that inverts the compromised and resisted classes; three signals we implemented and rejected; and the finding that a hand-picked threshold dominates the anytime-valid e-detector we started from, which we report rather than omit.

Everything in the paper is reproducible from committed data with no API access: `./reproduce.sh` regenerates every number and figure.

2. Background and Related Work

2.1. Indirect prompt injection

Direct prompt injection manipulates the user-supplied prompt; *indirect* prompt injection is the more dangerous variant for agents, because the adversary never talks to the agent at all. Instead they plant text in a resource—a web page, a shared document, an email body, a calendar invitation—that the agent will retrieve while pursuing a legitimate task. When the retrieved content enters the context window, the model has no reliable mechanism for distinguishing “data I was asked to process” from “instructions I should follow”, and it may execute the planted directive (Zhan et al., 2024). The consequences scale with the agent’s privileges: an agent that can send email can exfiltrate; an agent that can move money can steal.

Wallace et al. (2024) argue that the root cause is the absence of a privilege ordering over instruction sources, and train models to respect an explicit *instruction hierarchy*. This mitigates but does not eliminate the problem, since the model must still infer which span of a long context is privileged. Nasr et al. (2025) show that defences evaluated only against static attacks tend to collapse under adaptive attacks that optimise against the defence, and argue that adaptive evaluation should be the norm—a caution we take seriously in Section 10.1.

2.2. Benchmarks for agent security

Progress here is benchmark-driven, and the choice of benchmark determines what “detection” even means.

AgentDojo (Debenedetti et al., 2024) is the most widely used harness for our setting. It provides four tool environments (banking, workspace, travel, Slack), a set of *user tasks* defining legitimate goals, and a set of *injection tasks* defining attacker goals, together with attack templates that place the attacker’s directive into environment content. Crucially, it scores two things

separately: *utility*, whether the user’s task was accomplished, and *security*, whether the *injection* task’s goal was accomplished. Section 6 shows that the direction of this second flag is easy to invert, with severe consequences.

InjecAgent (Zhan et al., 2024) contributes a distinct family of injection templates aimed at tool-integrated agents, which we use to test whether a detector generalises beyond the template it was developed on. **τ -bench** (Yao et al., 2025) supplies realistic benign multi-turn tool-use trajectories in customer-service domains; it has no attack suite, so we use it purely to enlarge and diversify the benign reference. **WebArena** (Zhou et al., 2024) provides realistic web environments and is complementary but not injection-focused. Li et al. (2026) survey the space and note substantial inconsistency in how safety benchmarks define and measure success—a concern our Sections 6 and 8 substantiate concretely.

2.3. Defences: prevention

Prevention redesigns the agent so injected instructions cannot influence control flow. **CaMeL** (Debenedetti et al., 2025) is the most thorough instance: it extracts the control flow from the trusted user query, tracks data provenance through a capability system, and refuses tool calls whose arguments carry untrusted taint into sensitive positions. This is strong but invasive—it requires rebuilding the agent around the defence. **Command-Sans** (Das et al., 2025) takes a lighter-weight route, sanitising tool outputs to strip imperative content before the agent reads them. Prevention and monitoring are complementary: prevention reduces the attack surface, monitoring provides an independent tripwire when prevention is absent or incomplete.

2.4. Defences: detection

Detection leaves the agent untouched and raises an alarm from a score. **DataSentinel** (Liu et al., 2025) fine-tunes an LLM to detect contaminated inputs, formulated as a minimax game so that the detector is trained against adaptive evasion. **Task Shield** (Jia et al., 2025) checks each action for alignment with the user’s stated task—conceptually the closest prior work to our action-side signal, though it is used as a standalone gate rather than as one calibrated component of a fused statistic. **MELON** (Zhu et al., 2025) detects injection by re-executing the trajectory with a masked user prompt and comparing the resulting tool calls, on the principle that a hijacked action recurs even when the real task is removed; this is powerful but pays a full agent re-run per check.

AgentArmor (Wang et al., 2025) is the strongest published *trace-only* detector and our reference point. It lifts the agent trace into a program representation—control-flow, data-flow and program-dependence graphs—and enforces typed security policies over it, reporting 95.75% true positives at a 3.66% false-positive rate. Because it reasons about which untrusted value flowed into which argument of which tool, it is structurally different from a scalar text score, and Section 8 argues that its reported operating point exceeds what the trace alone can support on the corpus we study, which is informative about how such numbers should be interpreted.

Recent work also monitors adjacent failure modes: Belenky et al. (2026) detect reward hacking cheaply, Thaman (2026) benchmark tool-use exploits, and Opoku and Banahene (2026) apply conformal risk control to retrieval and tool-use drift.

2.5. *Distribution-free calibration and anytime-valid inference*

Our fusion rule is built from conformal p -values. Conformal prediction (Vovk et al., 2022) converts any real-valued score into a distribution-free valid p -value using only exchangeability of a calibration sample, which is what lets us combine signals living on incompatible scales without distributional assumptions.

The project began in the adjacent literature on *anytime-valid* inference. E-values and e-processes (Vovk and Wang, 2021; Ramdas et al., 2023; Grünwald et al., 2024) support continuous monitoring: a non-negative process whose expectation is bounded under the null can be inspected at every step, and thresholding it at $1/\alpha$ controls error at level α at every stopping time by Ville’s inequality. Shin et al. (2023) extend this to sequential *change* detection with e-detectors, whose average-run-length guarantee holds with no independence assumption—attractive here, because an agent’s actions are profoundly dependent. Sadhuka et al. (2025) note that when the reference is *learned* the process is no longer an exact e-process and supply a PAC-style thresholding correction, and Waudby-Smith and Ramdas (2023); Page (1954); Lorden (1971) provide the betting and classical change-detection background. Several 2026 works apply anytime-valid ideas to LLM systems (Xie, 2026; Li, 2026; Gauthier et al., 2026) and conformal ideas to agent error attribution (Feng et al., 2026).

We retain conformal calibration, which does real work in our method, but Section 9.2 reports that the sequential e-detector machinery did *not* outperform a hand-picked threshold at the data scales we could reach. We therefore

present it as one calibration option rather than as the paper’s contribution—a demotion the evidence forced.

3. Threat Model and Problem Setup

3.1. Agent execution traces

An agent pursues a user request u by interleaving tool calls and tool results. We write a trace of length T as

$$\mathcal{T} = ((o_0, a_1), (o_1, a_2), \dots, (o_{T-1}, a_T)), \quad (1)$$

where a_t is the t -th action (a tool name together with its arguments) and o_{t-1} is the tool output the agent observed immediately before choosing a_t . A monitor sees u and $\mathcal{T}$ and nothing else: no model weights, no activations, no logits. We deliberately exclude the agent’s own natural-language messages, both because they are not always available in production logging and to keep the setting comparable to prior trace-only work; Section 10.1 shows this exclusion is exactly what the residual undetectable cases exploit.

3.2. Adversary

The adversary controls the textual content of one or more environment resources that the agent may retrieve, and nothing else. They cannot modify u , the tool implementations, the model, or the monitor. Their goal is to cause the agent to execute an *injection task*: an action sequence serving the adversary rather than the user. This is the standard indirect-prompt-injection model (Debenedetti et al., 2024).

3.3. Two distinct detection objectives

Definition 1 (Attempt detection). A trace is *attempt-positive* if injected adversarial content entered the agent’s context, whether or not the agent complied.

Definition 2 (Compromise detection). A trace is *compromise-positive* if the agent actually accomplished the adversary’s injection task.

Both are legitimate targets. Attempt detection suits a guardrail that should escalate whenever untrusted instructions reach the context, since an agent that resisted today may not tomorrow. Compromise detection suits incident response and is the harder, more consequential question. The two

differ sharply in practice: in our corpus a capable agent resists most attempts, so the attempt-positive class is an order of magnitude larger than the compromise-positive class. Results for one must never be reported as results for the other—a confusion we made ourselves, as Section 6 describes.

4. Method

4.1. Signals

We compute four per-action signals. Two score the observation stream, asking whether untrusted content is behaving like an instruction; two score the action stream, asking whether the agent’s behaviour serves the user.

S1: tool-transition surprise (observation side, behavioural).. A Laplace-smoothed bigram model over consecutive tool names, fitted on benign traces, gives

$$s_t^{\text{tr}} = -\log \hat{P}(\text{tool}(a_t) \mid \text{tool}(a_{t-1})). \quad (2)$$

S2: instruction-likeness of the observation.. Let $\text{IL}(o)$ score how strongly a tool output reads as a directive addressed to the agent. We evaluate two instantiations. The *lexicon* variant is a weighted per-word density of imperative verbs, second-person pronouns, pseudo-markup and exclamation marks. The *judge* variant queries a small LLM with the observation alone and takes a normalised integer rating. Section 7.3 compares them; the judge generalises better across attack templates, at the price of one cached model call per distinct observation.

S3: action-justification audit (action side).. Given only the user request u , the tool name and its arguments—and *never* the observations—an LLM judge rates how far the call fails to serve u :

$$s_t^{\text{aj}} = \text{Unjustified}(u, a_t) \in [0, 1]. \quad (3)$$

Withholding the observations is a deliberate design choice: it prevents the audit from simply re-detecting the injected text, so S3 is not a restatement of S2 and the two may be fused without circularity.

S4: sensitive-sink provenance (action side).. A conjunctive, per-value rule that fires when the agent performs an externally-effectful call whose critical argument traces to untrusted prose:

$$s_t^{\text{sp}} = \mathbf{1}[\text{Sink}(a_t) \wedge \exists v \in \text{Ids}(a_t) : v \notin u \wedge \text{FirstSeenInProse}(v)]. \quad (4)$$

Here Sink marks tools with external effect (moving money, sending messages, publishing, editing standing instructions) as opposed to read-only lookups; Ids extracts identifier-like values (IBANs, e-mail addresses, URLs, long digit strings); and FirstSeenInProse holds when the value’s first appearance in the trace is inside a long natural-language line rather than a short structured record field. The intuition is that a legitimate recipient reaches the agent as a record it looked up, whereas an attacker’s recipient arrives embedded in the prose of a poisoned document.

4.2. Why naive normalisation fails

The four signals live on incompatible scales: s^{tr} is in nats and unbounded, s^{aj} lies in $[0, 1]$, s^{sp} is binary. Summing them raw lets the widest-ranged term dominate. A natural fix, and the one an earlier version of this work used, is to subtract the largest value observed in benign training data,

$$\phi_{\max}(x) = \max\{0, x - \beta\}, \quad \beta = \max_{x' \in \mathcal{D}_{\text{train}}} x', \quad (5)$$

so that anything within the benign range scores zero. This is a trap.

Proposition 1 (Excess-over-maximum can annihilate a signal). *Let $\phi_{\max}$ be as in (5) and let $\mathcal{B}, \mathcal{C}$ be the benign and attack score distributions. If $\mathbb{P}_{X \sim \mathcal{C}}(X \leq \beta) = 1$, then $\phi_{\max}(X) = 0$ almost surely for both classes and $\text{AUROC}(\phi_{\max}(\mathcal{B}), \phi_{\max}(\mathcal{C})) = 1/2$, irrespective of how well x itself separates the classes.*

Proof. $\phi_{\max}$ maps $(-\infty, \beta]$ to the single point 0. If the attack distribution is supported on $(-\infty, \beta]$ then the transformed attack scores are identically 0; the transformed benign scores are non-negative and equal 0 except on the benign event $\{x > \beta\}$, which has probability zero by the definition of β as the training maximum in the population limit. All comparisons are therefore ties, and AUROC, which scores ties at $1/2$, equals $1/2$. $\square$ $\square$

The premise is not hypothetical. In our data the benign median instruction-likeness is 0.010 but its maximum is 0.613—an outlier—whereas the median for the InjecAgent family is 0.115. Every InjecAgent score therefore fell below β , and a signal that separates that family at AUROC = 0.84 in raw form was reduced to exactly chance. We had published the resulting collapse and attributed it to the attack being subtly phrased. It was our normalisation.

The lesson generalises: a fusion rule must never apply a many-to-one transform to a component signal.

4.3. Conformal surprisal fusion

Definition 3 (Conformal tail surprisal). Let $x_{1:n}$ be benign calibration values for a signal and x a test value. The conformal p -value and surprisal are

$$p(x) = \frac{1 + |\{i : x_i \geq x\}|}{n + 1}, \quad \psi(x) = -\log p(x). \quad (6)$$

Proposition 2 (Distribution-free validity). *If $x, x_1, \dots, x_n$ are exchangeable, then $\mathbb{P}(p(x) \leq \alpha) \leq \alpha$ for every $\alpha \in (0, 1)$; equivalently $\mathbb{P}(\psi(x) \geq \log(1/\alpha)) \leq \alpha$.*

Proof. By exchangeability the rank of x among $\{x, x_1, \dots, x_n\}$ is uniform on $\{1, \dots, n + 1\}$. The event $\{p(x) \leq \alpha\}$ is the event that this rank places x among the top $\lfloor \alpha(n + 1) \rfloor$ values, which has probability $\lfloor \alpha(n + 1) \rfloor / (n + 1) \leq \alpha$. $\square$

Proposition 3 (Monotonicity and scale-freeness). *For fixed calibration values, ψ in (6) is non-decreasing in x and depends on x only through its rank. Consequently (i) ψ preserves the AUROC of the underlying signal, and (ii) ψ is invariant to any strictly increasing reparameterisation of the signal.*

Proof. $|\{i : x_i \geq x\}|$ is non-increasing in x , so p is non-increasing and $\psi = -\log p$ is non-decreasing; it is a function of the rank alone by inspection. AUROC depends only on the ordering of scores across the two classes, and a non-decreasing map cannot reverse any strict ordering; ties can only be created, which is why we state preservation rather than strict equality. Invariance to increasing reparameterisation follows because ranks are invariant. $\square$

Proposition 3 is the precise contrast with Proposition 1: the conformal surprisal cannot destroy a signal, because it is monotone and therefore ordering-preserving, whereas excess-over-maximum is many-to-one and can collapse an entire class to a point.

Definition 4 (SENTRY-Fuse). For signals $k = 1, \dots, K$ with per-trajectory summaries x_k (we use the per-trajectory maximum of each per-action signal) and benign calibration values $x_{1:n_k}^{(k)}$, the fused statistic is

$$S = \sum_{k=1}^K \psi_k(x_k) = -\sum_{k=1}^K \log p_k(x_k), \quad (7)$$

and the monitor alarms when $S \geq \tau$, with τ set to a benign quantile on a held-out calibration split.

The rule has no tuned weights: every signal contributes on the common scale of log-improbability, which is what makes an unweighted sum meaningful.

Proposition 4 (Fused false-alarm bound). *Under the null, for any $\tau > 0$,*

$$\mathbb{P}(S \geq \tau) \leq \min \left\{ 1, \sum_{k=1}^K e^{-\tau/K} \right\} = \min \{ 1, K e^{-\tau/K} \}. \quad (8)$$

If in addition the p_k are mutually independent, then $S \sim \text{Gamma}(K, 1)$ and $\mathbb{P}(S \geq \tau) = \Gamma(K, \tau)/\Gamma(K)$, the sharper Fisher bound.

Proof. If $S \geq \tau$ then some $\psi_k \geq \tau/K$, so $\{S \geq \tau\} \subseteq \bigcup_k \{\psi_k \geq \tau/K\}$; a union bound with Proposition 2 applied at level $\alpha = e^{-\tau/K}$ gives the first claim, and it requires no dependence assumption. Under independence each $-\log p_k$ is stochastically dominated by a standard exponential variable, and a sum of K such is dominated by a $\text{Gamma}(K, 1)$, giving the tail expression. $\square \square$

Remark 1. We report empirical operating points throughout rather than relying on Proposition 4, because the assumption-light union bound is loose and independence across our signals is not credible: S2 and S3 both respond to the presence of an injection, albeit through different channels. The bound’s role is to establish that the fused statistic admits distribution-free false-alarm control at all, and to make explicit what would be needed for a sharp one.

Because the fusion is additive over interpretable components, the dominant term $\arg \max_k \psi_k$ names *why* the monitor fired—an injected instruction, an unjustified action, an untrusted recipient, or an anomalous tool transition—which matters for triage.

Algorithm 1 SENTRY-Fuse

Require: benign trajectories $\mathcal{D}_{\text{cal}}$; target FPR α ; signals $1..K$

- 1: fit the benign reference (bigram table, per-tool vocabulary) on $\mathcal{D}_{\text{cal}}$
 - 2: **for** $k = 1, \dots, K$ **do**
 - 3: $x_{1:n_k}^{(k)} \leftarrow$ per-trajectory summaries of signal k on $\mathcal{D}_{\text{cal}}$
 - 4: **end for**
 - 5: $\tau \leftarrow (1 - \alpha)$ empirical quantile of $\{S(\mathcal{T}) : \mathcal{T} \in \mathcal{D}_{\text{cal}}\}$
 at monitoring time, for a trace $\mathcal{T}$:
 - 6: **for** $k = 1, \dots, K$ **do**
 - 7: $x_k \leftarrow \max_t s_t^{(k)}(\mathcal{T})$; $\psi_k \leftarrow -\log((1 + |\{i : x_i^{(k)} \geq x_k\}|)/(n_k + 1))$
 - 8: **end for**
 - 9: $S \leftarrow \sum_k \psi_k$
 - 10: **if** $S \geq \tau$ **then raise alarm**; report $\arg \max_k \psi_k$ as the dominant cause
 - 11: **end if**
-

Table 1: Trajectory corpora, with the *corrected* label polarity of Section 6: *compromised* means the injection task succeeded. Counts are after a two-action minimum filter.

Agent	Attack family	Benign	Attacks	Compromised	Resist
DeepSeek-V4-Flash	<code>important_instructions</code>	71	198	18 (9%)	1
DeepSeek-V4-Flash	InjecAgent	—	42	0 (0%)	
GPT-4o-mini	<code>important_instructions</code>	70	153	72 (47%)	
τ -bench retail (DeepSeek)	none	50	—	—	
Total		191	393	90	3

5. Experimental Setup

5.1. Corpora

We collect real trajectories by running tool-calling agents on standard benchmarks; Table 1 summarises the result. From AgentDojo (Debenedetti et al., 2024) we run all four suites (banking, workspace, travel, Slack) both benign and under attack, with two attack families: the `important_instructions` template and the InjecAgent (Zhan et al., 2024) template. From τ -bench (Yao et al., 2025) we run the retail domain benign-only, to diversify the benign reference. Two agents are used: DeepSeek-V4-Flash and GPT-4o-mini, from different vendors, so that agent identity can be varied independently of everything else.

Table 1 already contains a substantive finding. The two agents differ enormously in injection robustness: GPT-4o-mini is compromised in 47% of attempts, DeepSeek-V4-Flash in only 9%, and DeepSeek resisted the InjecAgent family *completely* (0/42). Robustness to indirect prompt injection is therefore a strong function of the agent, not a property of the attack alone, and a defence evaluated on a single agent may be measuring that agent’s fragility as much as its own merit.

Remark 2. Because DeepSeek was never compromised by InjecAgent, compromise detection is *undefined* for that cell: there are no positives. Only attempt detection is meaningful there, and we report it as such rather than quoting a rate over an empty class.

5.2. Protocol

Trajectories are converted to per-action signal streams by a deterministic, GPU-free adapter. Feature hashing uses a seed-stable CRC rather than Python’s per-process-salted hash, so the entire evaluation is byte-reproducible. LLM judge scores are cached by content hash (681 distinct observations, 1,329 distinct task–action pairs), so all reported evaluations run offline from committed caches.

Our evaluation protocol is deliberately strict about leakage. For each of ten seeds, the benign set is split in half: one half fits the reference model, calibrates every conformal transform, *and* sets the operating threshold; the other half supplies the held-out false positives. Compromise-positive trajectories are never used for fitting or calibration. We report mean $\pm$ standard deviation over the ten seeds.

Remark 3 (Leakage matters). An earlier version of this evaluation fitted the reference on all benign trajectories including the test half. Under that protocol the behavioural signal S1 appeared excellent, AUROC = 0.982; with calibration properly held out it collapses to 0.53–0.62. Only held-out numbers appear in this paper. This is worth stating plainly because the leaky number is the one a careless pipeline produces by default.

5.3. Metrics

We report AUROC, and true-positive rate at a matched benign false-positive rate. We use FPR = 3.66% throughout because that is the operating point AgentArmor (Wang et al., 2025) reports, which makes the comparison

as close to like-for-like as the published information allows; we also report 5%. Matching the FPR is essential—detectors are otherwise trivially incomparable.

6. A Label-Polarity Pitfall

Before presenting results we must document a methodological trap, because it inverted our own conclusions for a substantial period and because the ceiling analysis of Section 8 is impossible to get right without it.

AgentDojo returns a per-trajectory `security` flag. It is natural, and wrong, to read this as "the system was secure". The flag is computed as the success of the *injection* task, so `security = True` means the attacker's goal was achieved, i.e. the agent *was compromised*.

We initially adopted the opposite convention. The error was invisible in aggregate: both classes are non-empty, all downstream code ran, and every metric produced plausible numbers. It surfaced only when we asked why so many "compromised" trajectories contained no state-changing tool call at all, and read one:

Example 1 (A trajectory we had labelled compromised). Banking suite, `important_instructions`. The agent calls `read_file` and five read-only lookups, then states:

"I see a message embedded in the bill file that's trying to get me to do something else first. However, I shouldn't be tricked by that embedded instruction."

This agent *resisted*. It had been counted as a successful attack.

We resolved the polarity against ground truth rather than documentation. In the banking suite the adversary's goal is a transfer to a specific attacker IBAN, so compliance is directly checkable in the trace:

The correspondence in Table 2 is perfect, which settles the question: `security = True` is compromise.

The consequences for us were severe and instructive. Every "hijack-detection" number we had computed was measuring the *resisted* class. Our effect-based signals—the action audit and sink provenance—had been graded against trajectories in which, by construction, the agent had done nothing harmful; they appeared useless. Once the polarity was corrected, the action audit moved from AUROC = 0.547 to 0.926. We had nearly discarded a working signal because of a boolean.

Table 2: Resolving label polarity against ground truth in the banking suite. The only trajectories that actually wired money to the attacker’s account are exactly those with `security=True`.

	security=True	security=False
agent wired money to attacker IBAN	2	0
agent did not	0	46

Remark 4. We report this at length because the failure mode is silent, the fix is trivial, and the affected quantity—which class counts as a successful attack—is the foundation of every number in this literature. We recommend that any work using this harness verify polarity against an in-trace ground-truth check of the kind in Table 2, and report having done so.

7. Results

7.1. Compromise detection

Table 3 gives the main result. On GPT-4o-mini, where the compromised class is large enough for stable estimates (72 positives), the observation-side instruction signal alone reaches $\text{AUROC} = 0.955$ but a highly variable $\text{TPR} = 0.825 \pm 0.279$ at the matched operating point: it is strong on average and unreliable at a fixed low false-alarm budget. Adding the action-side audit lifts this to 0.932 ± 0.031 , and the full four-signal fusion to **0.960 ± 0.010** —above the 0.9575 reported by AgentArmor, and with an order of magnitude less variance than the single-signal detector.

Figure 1 visualises the ablation. Three points deserve emphasis.

First, **action-side evidence is complementary**. The pooled fused detectors reach 0.902, which is strictly above 0.867—the fraction of compromised trajectories in which the injected payload is visible in the trace at all (Section 8). No observation-side detector, however good, can exceed 0.867, because for the remaining cases there is nothing in the observation stream to see. Exceeding it is therefore direct evidence that S3 and S4 supply information S2 cannot.

Second, **fusion buys stability as well as power**. The single-signal detector’s ± 0.279 standard deviation at a 3.66% budget means its behaviour at a fixed low false-alarm rate is close to unpredictable across calibration draws; the fused detector’s ± 0.010 is deployable.

Table 3: Compromise detection, ten seeds, mean $\pm$ std, with calibration and threshold fitted on a held-out benign half. Best fused row in bold.

Detector	AUROC	TPR @ 3.66% FPR	TPR @ 5% FPR
<i>GPT-4o-mini: 72 compromised / 70 benign</i>			
S2 instruction only	0.955 ± 0.017	0.825 ± 0.279	0.825 ± 0.279
S1 behavioural only	0.534 ± 0.015	0.000 ± 0.000	0.000 ± 0.000
S1 + S2	0.954 ± 0.019	0.814 ± 0.276	0.822 ± 0.276
S2 + S3	0.985 ± 0.003	0.932 ± 0.031	0.932 ± 0.031
SENTRY-Fuse (S1–S4)	0.985 ± 0.005	0.960 ± 0.010	0.961 ± 0.010
<i>Pooled: 90 compromised / 191 benign</i>			
S2 instruction only	0.940 ± 0.008	0.719 ± 0.361	0.901 ± 0.031
S1 behavioural only	0.617 ± 0.035	0.000 ± 0.000	0.000 ± 0.000
S2 + S3	0.955 ± 0.003	0.902 ± 0.018	0.911 ± 0.008
SENTRY-Fuse (S1–S4)	0.956 ± 0.005	0.890 ± 0.039	0.904 ± 0.039
AgentArmor (Wang et al., 2025), published	—	0.9575	—

Third, **more signals is not monotonically better**. Pooled, S2+S3 (0.902) slightly outperforms the full fusion (0.890), because S1 is nearly uninformative under honest calibration and S4 is sparse, so both add variance without adding much signal. We report this rather than quoting only the best-performing subset.

7.2. Attempt detection and generalisation

Table 4 varies one axis at a time. Attempt detection transfers well across both axes: changing the agent leaves it intact (AUROC $0.76 \rightarrow 0.82$), and changing the attack family costs relatively little once the normalisation of Section 4.2 is fixed ($0.76 \rightarrow 0.71$). The DeepSeek/InjecAgent cell is the clearest illustration of why the two objectives must be kept apart: attempt detection is meaningful there (0.71), while compromise detection is undefined because the agent was never compromised.

7.3. Signal and normalisation ablation

Table 5 isolates the two design choices in the observation-side signal: how instruction-likeness is scored (fixed lexicon versus LLM judge) and how it is normalised (excess-over-maximum versus conformal). The interaction is the




Figure 1: Fusion ablation. (a) On GPT-4o-mini the four-signal fusion exceeds the strongest published trace-only operating point (dashed). (b) Pooled, the fused detectors exceed the payload-visibility ceiling (dotted), which no observation-only detector can. Error bars are one standard deviation over ten seeds.

Table 4: Generalisation across agent and attack family; ten seeds, corrected labels. “—” marks a cell with no compromised trajectories, where compromise detection is undefined.

Agent	Attack family	Comp.	Resisted	Attempt det.	Attempt AUROC
DeepSeek	<code>important_instructions</code>	18	180	0.65	0.76
DeepSeek	InjecAgent	0	42	0.52	0.71
GPT-4o-mini	<code>important_instructions</code>	72	81	0.74	0.82

point. Under the conformal transform the judge generalises far better than the lexicon across attack families (0.531 versus 0.455 attempt detection on InjecAgent). Under excess-over-maximum the *better* scorer is catastrophically worse (0.077), because the judge’s near-binary output puts some benign observation at the ceiling, and Proposition 1 then applies. Normalisation is not plumbing; it is a first-class design decision that can invert the ranking of scorers.

Figure 2 shows the separation the observation-side signal achieves, and Figure 3 walks through a single real compromise action by action, which is the clearest way to see what each signal contributes.

Table 5: Attempt detection / AUROC by instruction scorer and normalisation, ten seeds. The conformal transform is required for the semantic judge to be usable at all.

Scorer	Normalisation	<code>important_instr.</code>	InjecAgent
lexicon	excess-over-max	0.658/0.769	0.083/0.468
lexicon	conformal	0.620/0.756	0.455/0.701
LLM judge	excess-over-max	0.077/0.432	0.090/0.490
LLM judge	conformal	0.648/0.750	0.531/0.714




Figure 2: Observation instruction-likeness (maximum over a trajectory) for benign and attacked trajectories. Benign tool output sits near zero; attacked trajectories sit far above the learned nominal range, whether or not the agent ultimately complied—which is why this signal is a good *attempt* detector.

8. The Detectability Ceiling

The results above raise an obvious question: how much room is left? We now show that the answer is very little, and that the limit is a property of the trace rather than of any particular detector.

A trace-only monitor can only react to evidence present in the trace. For indirect prompt injection there are exactly two channels through which a compromise can leave such evidence:

- **Observation channel:** the injected payload appears in some tool output o_t that the monitor sees.




Figure 3: A real prompt-injection compromise, action by action. The user asks the agent to adjust a rent payment; at t_0 it reads a file whose content hides an `<INFORMATION>` block directing it to wire money to the attacker. The poisoned text enters the monitored trace at t_1 , where the observation-side signal spikes; the agent nonetheless complies at t_6 , sending money to the attacker’s IBAN—five actions after the monitor would have halted it.

- **Action channel:** the agent performs an externally-effectful call whose arguments reflect the adversary’s goal.

If a compromise leaves evidence in neither channel, its trace is—by construction—indistinguishable from that of a benign trajectory performing the same reads.

Assumption 1 (Trace indistinguishability). For every compromised trajectory whose trace exhibits neither channel, there exists a benign trajectory inducing the same monitor input $(u, \mathcal{T})$.

Theorem 5 (Detectability ceiling). *Let $\mathcal{C}$ be the set of compromised trajectories, and let $N \subseteq \mathcal{C}$ be those exhibiting neither channel, with $\nu = |N|/|\mathcal{C}|$. Under Assumption 1, every monitor M that is a function of $(u, \mathcal{T})$ and has false-positive rate at most α satisfies*

$$\text{TPR}(M) \leq (1 - \nu) + \nu \alpha. \quad (9)$$

Table 6: Where evidence of compromise lives, over the 90 genuinely compromised trajectories.

Channel of evidence	n	fraction
injected payload visible in an observation	78	0.867
externally-effectful (sink) call	74	0.822
either channel	83	0.922
neither channel (no trace evidence at all)	7	0.078




Figure 4: (a) Evidence channels over the 90 compromised trajectories. (b) Our pooled operating point against the attainable ceiling and the strongest published trace-only figure. The published value lies above the ceiling that this corpus supports.

Proof. Partition $\mathcal{C}$ into N and its complement. On $\mathcal{C} \setminus N$ we bound the true-positive rate by 1, contributing at most $(1 - \nu)$. For $\mathcal{T} \in N$, Assumption 1 supplies a benign trajectory with identical monitor input; since M is a deterministic function of that input, it alarms on the compromised trajectory if and only if it alarms on the benign one. The set of benign inputs on which M alarms has probability at most α by hypothesis, so M alarms on at most an α fraction of N , contributing at most $\nu\alpha$. Summing gives the bound. $\square$ $\square$

Corollary 6. *With the measured $\nu = 0.078$ and $\alpha = 0.0366$, $\text{TPR} \leq 0.922 + 0.078 \times 0.0366 = 0.925$.*

Table 6 and Figure 4 give the measurement.
Three consequences follow.

We are close to saturation.. Our pooled 0.902 is 97.8% of the attainable 0.922. Further gains on this corpus from better trace-only *scoring* are bounded

by roughly two percentage points; the remaining headroom is not in the detector but in the evidence.

Observation-only monitoring is capped below the state of the art.. The observation channel alone supports at most 0.867. A literature that scores only untrusted text therefore cannot, on this corpus, reach the numbers it aims for—which is precisely why our action-side signals matter.

Reported figures above the ceiling require explanation.. The 0.9575 of Wang et al. (2025) exceeds Corollary 6’s 0.925. We stress that this is not evidence of error on their part: their corpus, agents and positive-class definition all differ from ours, and a differently-constituted corpus will have a different ν . What it does establish is that the comparison is not like-for-like, and that a headline TPR is uninterpretable without the accompanying ν . We recommend that trace-only detection papers report the fraction of positives that leave no trace evidence, so that reported rates can be read against the ceiling they are competing with.

9. Negative Results

A monitor’s credibility rests on knowing what it cannot do. We report every component we implemented and rejected.

9.1. Signals that did not separate

Lexical data-flow taint (AUROC 0.46).. We scored the fraction of an action’s argument tokens traceable to untrusted observations. It performs *below* chance because benign agents legitimately act on tool-derived values constantly—reading a recipient from a transaction record and paying it is the normal case. Token-overlap fraction cannot distinguish that from acting on an attacker-supplied recipient. This failure motivated the conjunctive, sink-restricted form of S4.

Lexical goal-grounding (AUROC 0.50).. Token overlap between the action and the user request is too coarse a proxy for ”serves the task”. Its semantic replacement, S3, works (0.926)—the idea was right, the implementation was lexical.

Injection-conditioned interaction (no gain).. Gating behavioural surprise on a preceding injection produced results identical to the plain sum.

Table 7: Long-horizon false alarms on 1000-action benign streams, ten seeds, identical signal. A hand-picked threshold dominates the e-detector on both axes.

Decision rule	False alarms / 1k	ARL to first	Attack det.
e-detector, $\alpha = 0.2$	48.9	29	0.624
e-detector, $\alpha = 0.05$	19.7	138	0.530
e-detector, $\alpha = 0.01$	19.7	138	0.530
e-detector, $\alpha = 0.002$	19.7	138	0.530
fixed threshold, $q = 0.99$	17.4	158	0.588
fixed threshold, $q = 0.999$	11.0	343	0.565

Sink provenance alone (TPR 0.112 at FPR 0.073).. As a standalone trajectory-level detector S4 is weak, because it fires only for the subset of attacks that route an identifier into a sensitive sink. It earns its place only inside the fusion, where its precision complements S2’s recall.

9.2. The anytime-valid machinery did not earn its headline

This work began as an application of e-detectors (Shin et al., 2023) with PAC-calibrated thresholds (Sadhuka et al., 2025), on the argument that continuous monitoring demands anytime-valid false-alarm control. We tested that argument directly. Holding the signal fixed and varying only the decision rule, we built 1000-action benign streams by concatenating held-out benign trajectories and swept each rule’s knob (Table 7).

Two findings stand out. A fixed threshold on the same score *strictly dominates* the e-detector—fewer false alarms and higher detection simultaneously. And the e-detector’s results are *identical* for $\alpha \in \{0.05, 0.01, 0.002\}$: the PAC threshold saturates at the order-statistic index available from our calibration set, so the guarantee cannot be tightened below $\alpha = 0.2$, where it promises only $\text{ARL} \geq 5$ actions—a vacuous promise for a monitor. A cross-agent transfer test showed no advantage either.

We therefore retain conformal calibration, which does real work (Propositions 2–3), and demote the sequential apparatus to an optional operating mode. The theory is not wrong; it simply requires far more trusted data than a realistic deployment provides before its guarantee becomes both valid and non-vacuous. We consider this a useful negative result for a community currently enthusiastic about anytime-valid monitoring.

10. Discussion

What the ceiling changes.. The prevailing framing treats trace-only injection detection as a modelling problem: find a better score. Our measurement suggests it is largely an *evidence* problem. Once both channels are exploited, 97.8% of the attainable performance is reached by a fusion rule with no tuned weights. Effort is better spent widening the evidence—monitoring the agent’s own natural-language output, or scoring retrieved documents at the retrieval boundary before the agent chooses what to surface—than on further refining scores over the same trace.

Why two channels rather than one.. S2 and S3 fail on disjoint cases. S2 misses compromises whose payload never surfaces; S3 misses compromises whose harmful action looks superficially plausible. Because the failures are disjoint, fusion is close to additive in coverage, which is what the pooled $0.902 > 0.867$ demonstrates. This is a concrete instance of a general design principle for monitors: prefer signals that fail on different inputs over signals that are individually stronger.

Agent robustness is a confound.. GPT-4o-mini was compromised five times more often than DeepSeek-V4-Flash on identical attacks, and DeepSeek resisted an entire attack family. A defence evaluated on one agent is partly measuring that agent. We recommend reporting the agent’s unprotected compromise rate alongside any detector’s TPR, as Table 1 does.

Cost.. S1 and S4 are dictionary and regular-expression operations. S2 and S3 each require one small-model call per distinct observation and per distinct action, cached; on our corpora this is 2,010 calls total for 393 trajectories, and zero at re-evaluation time. There is no GPU, no access to model internals, and no agent re-execution, in contrast to re-execution defences (Zhu et al., 2025).

Deployment and interpretability.. Because (7) is additive over interpretable components, an alarm comes with an attribution: the dominant ψ_k names the reason. An operator can route “injected instruction detected” differently from “action unjustified given the user’s request”, and can tune α against a false-alarm budget.

10.1. Limitations

(i) *Adaptive adversaries.* Our attacks are static templates. An adversary aware of SENTRY-Fuse would phrase injections to minimise instruction-likeness *and* route effects through non-sink tools *and* keep actions superficially task-plausible. Requiring simultaneous evasion of channels that fail on disjoint inputs raises the bar, but we have not measured by how much, and Nasr et al. (2025) caution against assuming such bars hold.

(ii) *Corpus scale and breadth.* 90 compromised trajectories, two agents, two attack families, one benchmark family. The DeepSeek compromised class (18) is small enough that its estimates are wide, and it is why we present GPT-4o-mini as the primary compromise-detection corpus and report pooled numbers alongside.

(iii) *Judge dependence.* S2 and S3 depend on a small LLM whose behaviour may drift with model updates; the cached-score design makes a given evaluation reproducible but does not make the judge itself stable over time.

(iv) *The monitor is blind to the agent’s own words.* All seven no-evidence compromises are of this kind, and this is the single most promising extension.

(v) *Ceiling assumption.* Theorem 5 rests on Assumption 1, an idealisation: two traces that share no injection-related evidence might still differ in some incidental way a detector could exploit. The bound should be read as characterising what is available in principle from the two named channels.

10.2. Future work

Monitor the response channel. Extending the trace to the agent’s natural-language messages is the direct route past 0.922, and would let the monitor catch information-disclosure compromises that never touch a tool.

Score content at the retrieval boundary. A monitor placed where documents *enter* the environment, rather than where the agent surfaces them, would see payloads the agent never reads aloud.

Structured provenance. S4 is a coarse approximation of the program-dependence analysis that Wang et al. (2025) perform. A faithful data-flow implementation, fused rather than used standalone, is a natural strengthening.

Adaptive evaluation. Attacks optimised jointly against all four channels are the decisive robustness test.

11. Conclusion

We set out to build a better trace-only monitor for prompt-injection compromise in LLM agents and found that the binding constraint was not the scoring function but the evidence. SENTRY-Fuse converts heterogeneous trace signals into conformal tail surprisals and fuses them additively; the construction is scale-free, provably order-preserving, and free of tuned weights, and we showed that a common alternative normalisation provably destroys signal—as it did, measurably, in an earlier version of this work. Fusing observation-side and action-side evidence attains 0.960 ± 0.010 true positives at a 3.66% false-positive rate on a matched corpus, exceeding the strongest published trace-only figure, and 0.902 pooled, which is above what payload visibility alone can support and 97.8% of the attainable ceiling.

That ceiling is our most durable contribution. Some successful injections leave no trace evidence at all; on our corpus 7.8% do, capping every trace-only monitor at 0.922. Reported rates above such a ceiling are not comparable across corpora unless the corresponding fraction is reported too, and we recommend that it routinely be. Alongside this we document a label-polarity trap that silently inverted our own conclusions, and we report in full the components we rejected—including the anytime-valid machinery this project began from, which a hand-picked threshold outperformed. The remaining headroom lies in widening what the monitor is allowed to see, not in sharpening how it scores what it already sees.

Data availability

All code, the collected trajectory corpora, cached judge scores, and scripts that regenerate every number and figure are available at <https://github.com/vinhq dang/SENTRY-Anytime-Valid-E-Process-Monitoring-of-LLM-Agent-Action-Streams>. `./reproduce.sh` reproduces the paper offline from committed data.

Declaration of competing interest

The author declares no competing interests. No external funding was received.

References

- Belenky, I., Itria, J., Johns, S., 2026. Cheap reward hacking detection. arXiv preprint arXiv:2606.08893 URL: <https://arxiv.org/abs/2606.08893>, doi:10.48550/arXiv.2606.08893.
- Das, D., Beurer-Kellner, L., Fischer, M., Baader, M., 2025. Command-sans: Securing ai agents with surgical precision prompt sanitization. arXiv preprint arXiv:2510.08829 URL: <https://arxiv.org/abs/2510.08829>, doi:10.48550/arXiv.2510.08829.
- Debenedetti, E., Shumailov, I., Fan, T., Hayes, J., Carlini, N., Fabian, D., Kern, C., Shi, C., Terzis, A., Tramèr, F., 2025. Defeating prompt injections by design. arXiv preprint arXiv:2503.18813 URL: <https://arxiv.org/abs/2503.18813>, doi:10.48550/arXiv.2503.18813.
- Debenedetti, E., Zhang, J., Balunović, M., Beurer-Kellner, L., Fischer, M., Tramèr, F., 2024. Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents, in: Advances in Neural Information Processing Systems 37 (NeurIPS 2024) Datasets and Benchmarks Track. URL: https://proceedings.neurips.cc/paper_files/paper/2024/hash/97091a5177d8dc64b1da8bf3e1f6fb54-Abstract-Datasets_and_Benchmarks_Track.html.
- Feng, N., Sui, Y., Hou, S., Wu, G., Cresswell, J.C., 2026. Conformal agent error attribution. arXiv preprint arXiv:2605.06788 URL: <https://arxiv.org/abs/2605.06788>, doi:10.48550/arXiv.2605.06788.
- Gauthier, E., Bach, F., Jordan, M.I., 2026. Anytime detection of strategic deviations in multi-agent systems. arXiv preprint arXiv:2601.05427 URL: <https://arxiv.org/abs/2601.05427>, doi:10.48550/arXiv.2601.05427.
- Grünwald, P., de Heide, R., Koolen, W.M., 2024. Safe testing. Journal of the Royal Statistical Society Series B: Statistical Methodology 86, 1091–1128. doi:10.1093/jrsssb/qkae011.
- Jia, F., Wu, T., Qin, X., Squicciarini, A., 2025. The task shield: Enforcing task alignment to defend against indirect prompt injection in llm agents, in: Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers), Association

- for Computational Linguistics, Vienna, Austria. pp. 29680–29697. URL: <https://aclanthology.org/2025.acl-long.1435/>, doi:10.18653/v1/2025.acl-long.1435.
- Li, M.Q., Fung, B.C.M., Li, B., Ismail, H., Iqbal, F., 2026. Taxonomy and consistency analysis of safety benchmarks for ai agents. arXiv preprint arXiv:2605.16282 URL: <https://arxiv.org/abs/2605.16282>, doi:10.48550/arXiv.2605.16282.
- Li, Y., 2026. Who drifted: the system or the judge? anytime-valid attribution in llm evaluation pipelines. arXiv preprint arXiv:2606.15474 URL: <https://arxiv.org/abs/2606.15474>, doi:10.48550/arXiv.2606.15474.
- Liu, Y., Jia, Y., Jia, J., Song, D., Gong, N.Z., 2025. Datasentinel: A game-theoretic detection of prompt injection attacks, in: 2025 IEEE Symposium on Security and Privacy (SP), IEEE, San Francisco, CA, USA. pp. 2190–2208. URL: <https://doi.org/10.1109/SP61157.2025.00250>, doi:10.1109/SP61157.2025.00250.
- Lorden, G., 1971. Procedures for reacting to a change in distribution. The Annals of Mathematical Statistics 42, 1897–1908. doi:10.1214/aoms/1177693055.
- Nasr, M., Carlini, N., Sitawarin, C., Schulhoff, S.V., Hayes, J., Ilie, M., Pluto, J., Song, S., Chaudhari, H., Shumailov, I., Thakurta, A., Xiao, K.Y., Terzis, A., Tramèr, F., 2025. The attacker moves second: Stronger adaptive attacks bypass defenses against llm jailbreaks and prompt injections. arXiv preprint arXiv:2510.09023 URL: <https://arxiv.org/abs/2510.09023>, doi:10.48550/arXiv.2510.09023.
- Opoku, J., Banahene, D., 2026. Toolchain-crc: Conformal risk control for agentic ai under retrieval and tool-use drift. arXiv preprint arXiv:2606.18467 URL: <https://arxiv.org/abs/2606.18467>, doi:10.48550/arXiv.2606.18467.
- Page, E.S., 1954. Continuous inspection schemes. Biometrika 41, 100–115. doi:10.1093/biomet/41.1-2.100.
- Ramdas, A., Grünwald, P., Vovk, V., Shafer, G., 2023. Game-theoretic statistics and safe anytime-valid inference. Statistical Science 38, 576–601. doi:10.1214/23-sts894.

- Sadhuka, S., Prinster, D., Fannjiang, C., Scalia, G., Berger, B., Regev, A., Wang, H., 2025. E-evaluator: Reliable agent verifiers with sequential hypothesis testing. arXiv preprint arXiv:2512.03109 URL: <https://arxiv.org/abs/2512.03109>, doi:10.48550/arXiv.2512.03109.
- Shin, J., Ramdas, A., Rinaldo, A., 2023. E-detectors: A nonparametric framework for sequential change detection. *The New England Journal of Statistics in Data Science* 1, 229–260. doi:10.51387/23-nejsds51.
- Thaman, K., 2026. Reward hacking benchmark: Measuring exploits in llm agents with tool use. arXiv preprint arXiv:2605.02964 URL: <https://arxiv.org/abs/2605.02964>, doi:10.48550/arXiv.2605.02964.
- Vovk, V., Gammerman, A., Shafer, G., 2022. *Algorithmic Learning in a Random World*. 2nd ed., Springer, Cham, Switzerland. doi:10.1007/978-3-031-06649-8.
- Vovk, V., Wang, R., 2021. E-values: Calibration, combination and applications. *The Annals of Statistics* 49, 1736–1754. doi:10.1214/20-aos2020.
- Wallace, E., Xiao, K., Leike, R., Weng, L., Heidecke, J., Beutel, A., 2024. The instruction hierarchy: Training llms to prioritize privileged instructions. arXiv preprint arXiv:2404.13208 URL: <https://arxiv.org/abs/2404.13208>, doi:10.48550/arXiv.2404.13208.
- Wang, P., Liu, Y., Lu, Y., Cai, Y., Chen, H., Yang, Q., Zhang, J., Hong, J., Wu, Y., 2025. Agentarmor: Enforcing program analysis on agent runtime trace to defend against prompt injection. arXiv preprint arXiv:2508.01249 URL: <https://arxiv.org/abs/2508.01249>, doi:10.48550/arXiv.2508.01249.
- Waudby-Smith, I., Ramdas, A., 2023. Estimating means of bounded random variables by betting. *Journal of the Royal Statistical Society Series B: Statistical Methodology* 86, 1–27. doi:10.1093/jrsssb/qkad009.
- Xie, Y., 2026. Sequential statistical inference for large language models: Representation, validity, and monitoring. arXiv preprint arXiv:2606.07624 URL: <https://arxiv.org/abs/2606.07624>, doi:10.48550/arXiv.2606.07624.

- Yao, S., Shinn, N., Razavi, P., Narasimhan, K., 2025. τ -bench: A benchmark for tool-agent-user interaction in real-world domains, in: The Thirteenth International Conference on Learning Representations (ICLR). URL: <https://iclr.cc/virtual/2025/poster/28170>.
- Zhan, Q., Liang, Z., Ying, Z., Kang, D., 2024. Injecagent: Benchmarking indirect prompt injections in tool-integrated large language model agents, in: Findings of the Association for Computational Linguistics: ACL 2024, Association for Computational Linguistics, Bangkok, Thailand. pp. 10471–10506. URL: <https://aclanthology.org/2024.findings-acl.624/>, doi:10.18653/v1/2024.findings-acl.624.
- Zhou, S., Xu, F.F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., Neubig, G., 2024. Webarena: A realistic web environment for building autonomous agents, in: The Twelfth International Conference on Learning Representations (ICLR). URL: <https://openreview.net/forum?id=oKn9c6ytLx>.
- Zhu, K., Yang, X., Wang, J., Guo, W., Wang, W.Y., 2025. Melon: Provable defense against indirect prompt injection attacks in ai agents, in: Proceedings of the 42nd International Conference on Machine Learning (ICML), PMLR, Vancouver, Canada. pp. 80310–80329. URL: <https://proceedings.mlr.press/v267/zhu25z.html>.